

1 The Memorist Tale: Once Annotated Lazy Program Cost is Fun

2
3 ANONYMOUS AUTHOR(S)

4 Lazy evaluation offers great flexibility in computing only what is necessary. However, reasoning about the
5 cost of lazy programs is notoriously difficult because computation happens out of order and depends on
6 future demands. Recently, researchers are looking into alternative semantics that models the cost of lazy
7 evaluation that enables simple cost analysis without reasoning about states. However, existing approaches
8 either require dealing with nondeterminism, or require a complex bidirectional semantics. In this paper, we
9 introduce the *Memorist Semantics*, a novel semantics to analyse the cost of lazy programs by tracking the
10 cost and dependency of computing each component. Our approach annotates each component of a term with
11 fine-grained cost information and usage sets, resulting in a deterministic semantics that can be encapsulated
12 with a simple monadic interface. We formalize our semantics in the Rocq Prover and verify its soundness with
13 respect to the existing Clairvoyance Semantics.

14 Additional Key Words and Phrases: computation cost, lazy evaluation, operational semantics
15

16 ACM Reference Format:

17 Anonymous Author(s). 2025. The Memorist Tale: Once Annotated Lazy Program Cost is Fun. In *Proceedings of*
18 *ESOP*. ACM, New York, NY, USA, 31 pages. <https://doi.org/XX.XXXX/XXXXXXX>

19 20 Introduction

21 A key feature of pure functional programming languages is referential transparency: pure functions
22 yield identical outputs for identical inputs. This facilitates formally verifying functional correctness,
23 as one can reason equationally about *what* the function does in *isolation*; without considering
24 external state or any other external information.

25 However, since functional programs typically abstract away *how* programs evaluate, assessing
26 the computational cost of a function presents challenges. In particular, how fast a function evaluates,
27 does not only depend on the function definition but it is intricately tied to the evaluation strategy
28 being used. Two common strategies used in functional programming languages are: *call-by-value*
29 which is used in eager programming languages such as Standard ML, and *call-by-need* which is
30 employed by lazy programming languages, such as Haskell.

31 The call-by-need evaluation strategy enables expressive and flexible programming by avoiding
32 redundant or unnecessary computation. Unlike eager evaluation, lazy evaluation ensures that
33 components of a term are evaluated at most once, and only if they are required. This property
34 allows for elegant program structuring and efficiency improvements, but it comes at a cost: analysing
35 the cost of lazy programs remains notoriously difficult.

36
37 *Running Example.* We consider the following simple example to explain the difference Let
38 `truePrefix` be a function that takes a list of booleans as input and returns the prefix of the input
39 list that only contains true elements. The function `truePrefixOfAppend`, that we use as a running
40 example in this paper, takes two lists of booleans, appends them, and applies `truePrefix` on the

41
42 Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee
43 provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and
44 the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses,
45 contact the owner/author(s).

46 *ESOP*,

47 © 2025 Copyright held by the owner/author(s).

48 ACM ISBN 978-x-xxxx-xxxx-x/YY/MM

49 <https://doi.org/XX.XXXX/XXXXXXX>

```

50   Fixpoint truePrefix (xs : list bool) : list bool :=
51     match xs with
52       | nil => nil
53       | cons x xs' =>
54         let zs := truePrefix xs' in
55           if x then cons x zs else nil
56       end.
57
58   Definition truePrefixOfAppend (xs ys : list bool) : list bool := let zs := append xs ys in truePrefix zs.
59
60
61 Fig. 1. The Gallina definitions of append, truePrefix, and truePrefixOfAppend placed in A-normal form.
62
63
64 appended list. Figure 1 presents an implementation of these functions in Gallina; the specification language of the Rocq Prover. Now consider applying this example to two concrete lists: truePrefixOfAppend [true;false] [true].
65 We are interested in the cost in terms of time this program takes to evaluate. A straightforward way of calculating such a cost is by counting the number of function applications. While Gallina does not enforce a particular evaluation strategy, we can imagine that this code is a shallow embedding of a program in another language. For example, one can use hs-to-coq [Breitner et al. 2021; Spector-Zabusky et al. 2018] to automatically translate Haskell code to this form. The evaluation cost depends on the evaluation strategy used in the source language.
66
67
68
69
70
71 Eager vs. lazy evaluation. Under eager evaluation, the computation of truePrefixOfAppend first computes append of the two lists [true; false] and [true] in full, evoking 3 calls to append. Then the resulting value appending the two lists is passed to truePrefix. Every element is processed, thus evoking 4 calls to truePrefix (including one for the empty list). Overall, the program evaluates to [true] and incurs a cost of  $3+4+1 = 8$ , with one additional call to the top-level truePrefixOfAppend.
72
73
74
75
76 In comparison, lazy evaluation is demand-driven: computation is performed based on which parts of the output are demanded. Let's start by assuming that the demand is only the weak-head normal form (WHNF) of the list.
77
78
79 Figure 2 illustrates how the evaluation proceeds stepwise. In the case of lazy evaluation, we only need the result up to cons true t, with t pointing to a thunk that need not be evaluated. Only one step of truePrefix, and one step of append is needed. Adding an additional step due to the call to truePrefixOfAppend itself, the total computation cost adds up to 3. At this step, we have all we need for WHNF, so the evaluation stops at this step marked (*).
80
81
82
83
84 However, evaluating to WHNF may not be the end of the story in every context—more of the list could be demanded in the future. For example, let's assume that another computation that happens later might demand the full list from truePrefixOfAppend. Fortunately, we have already computed such list cons true t, so we can resume from there instead of starting from the beginning—this is known as sharing in lazy evaluation. In this case, t must be evaluated as well; that is, we need the result to be fully reduced to cons true nil. One more step of truePrefix and of append is needed, making the computation cost 5 in total.
85
86
87
88
89
90
91 Even in this case, the cost is smaller than that incurred by eager evaluation, as computations not necessary for producing the final result are never performed. This includes, for instance, the expression append [] [true] that is bound to  $z_2$  in the illustration.
92
93
94
95 Challenges. The example demonstrates several challenges in analysing the cost of lazy evaluation:
96
97
98

```

- *The cost of lazy evaluation is not local.* Unlike eager evaluation, the cost incurred by a function may not happen at the call site of the function—it might happen later inside

```

99      truePrefixOfAppend (true :: false :: []) (true :: [])
100     = ✓let zs := append (true :: false :: []) (true :: []) in truePrefix zs
101     = let zs := ✓(let z1 := append (false :: []) (true :: []) in true :: z1) in truePrefix zs
102     = let z1 := append (false :: []) (true :: []) in truePrefix (true :: z1)
103     = let z1 := ... in ✓let w1 := truePrefix z1 in true ↗ true :: w1
104     = let z1 := ... in let w1 := truePrefix z1 in true :: w1
105     = let z1 := ✓(let z2 := append [] (true :: []) in false :: z2) in
106       let w1 := truePrefix z1 in true :: w1
107     = let z2 := append [] (true :: []) in let w1 := truePrefix (false :: z2) in true :: w1
108     = let z2 := ... in let w1 := ✓(let w2 := truePrefix z2 in false ↗ false :: w2) in true :: w1
109     = let z2 := ... in let w1 := (let w2 := ... in []) in true :: w1
110     = let z2 := ... in true :: []
111

```

Fig. 2. A stepwise illustration of lazily evaluating the program `truePrefixOfAppend`. Here we choose to count function applications. A tick (\checkmark) is placed immediately after unfolding a step of function application to indicate that a cost is incurred.

another function's application. The cost of a function can depend on future demand. This makes it challenging to analyse individual lazy functions in isolation.

- *Evaluation steps of different functions are interleaved.* As demonstrated in Figure 2, the evaluation steps of `truePrefix` and `append` are interleaved to evaluate every element in the result list of `truePrefixOfAppend`.
- *The evaluation is stateful.* This is especially important for modelling sharing, as the evaluation needs to remember what has been previously computed, and what has not.

Existing approaches. How can we deal with these challenges and analyse lazy computation cost then? We can treat any lazy programs as a stateful program and analyse them using program logics for reasoning about states, such as separation logic. Pottier et al. [2024] took this approach utilizing the Iris^{\$} framework [Mével et al. 2019] and the Iris separation logic [Spies et al. 2022]. But what if we just reason about lazy functional programs using simple tools like *equational reasoning*? Danielsson [2008] proposed a simple approach based on a graded monad that tracks the computation cost. The approach was later adopted by Handley et al. [2020] to be used in Liquid Haskell [Vazou 2016]. However, this approach requires the user to know future demand to correctly deal with sharing.

Instead of directly dealing with the stateful natural semantics of laziness [Launchbury 1993], another approach is to use an *alternative semantics* that is equivalent to lazy evaluation in terms of computational cost. The key idea is that it only matters *whether* a computation happens, not *when* the computation happens, when analyzing computational cost. This makes it possible to *localise* lazy computation cost, provided there is a way to know future demands. The Clairvoyance Semantics [Hackett and Hutton 2019] takes this approach and it simulates future demands via *nondeterminism*. Li et al. [2021] further show that this semantics can be encoded in a simple interface using the clairvoyance monad. However, nondeterminism makes both testing and formal reasoning challenging. Indeed, Li et al. [2021] proposed an additional logic similar to Incorrectness Logic [O'Hearn 2020] for this reason.

To avoid nondeterminism, Xia et al. [2024] proposed the Demand Semantics based on the prior work of Bjerner and Holmström [1989]. The Demand Semantics employs two evaluation directions: a *forward* evaluation that, given input, computes output as normal; and a *demand* evaluation that, given pure input and an *output demand*, calculates the *minimal input demand* and the corresponding computation cost. This approach was shown to be effective in formal reasoning, but the two

148 evaluation directions mean that one function needs to be translated into two different functions for
 149 formal reasoning, making the translation error-prone and bringing in new challenges to formal
 150 reasoning. We defer a more detailed discussion contrasting these approaches to [Section 6](#).

151 *Our key idea.* In this paper, we propose the Memorist Semantics, a novel cost semantics for lazy
 152 evaluation that tracks both usage and cost of terms. The key idea of the Memorist Semantics is to
 153 run an eager evaluation, while giving every piece of data a unique name and “memorising” which
 154 other information was used for computing this data. The Memorist Semantics enables analysing
 155 computation cost *locally* in isolation without considering states, similar to the Clairvoyance Se-
 156 mantics and the Demand Semantics. In the meantime, the Memorist Semantics improves on the
 157 Clairvoyance Semantics in that it is *deterministic* and it improves on the Demand Semantics, in
 158 that it only employs one semantics and does not require code duplication.
 159

160 To achieve this, we address two key challenges: The first challenge is *precise cost attribution* for
 161 different components of a term or value. Since lazy evaluation may only require parts of a term,
 162 we must ensure that unnecessary computation is excluded from cost calculations. Our approach
 163 achieves this by tracking the cost of computing each component of a value separately, annotating
 164 each component of a data structure with its cost. The second challenge is *accounting for shared*
 165 *computation*. Because lazy evaluation ensures that each component of a term is evaluated at most
 166 once, even if its value is used multiple times, our semantics must prevent duplicate cost accounting.
 167 We achieve this by tracking *usage sets* and using set union for bookkeeping, ensuring that each
 168 component’s cost is included at most once, regardless of how often the value is reused.

169 *Contributions.* We make the following contributions:

- 170 • We introduce the Memorist Semantics, a deterministic cost semantics for call-by-need that
 171 tracks both the cost and usage of subexpressions ([Section 3](#)).
- 172 • A proof of the correctness of our semantics, as well as a proof that our cost model is sound
 173 and that it correctly predicts the execution cost of a program under call-by-need ([Section 4](#)).
- 174 • We show that the Memorist Semantics can be encoded using a simple monadic interface just
 175 like the Clairvoyance Semantics via a proof of concept implementation in Gallina ([Section 5](#)).
- 176 • A formalization of our results in the Rocq Prover (formerly Coq) [[Coq development team](#)
 177 [2024](#)], ensuring rigorous correctness proofs.

178 We describe the intuition behind the Memorist Semantics in [Section 2](#) and discuss related work
 179 in [Section 6](#). We conclude the paper and discuss future work in [Section 8](#).
 180

181 **2 The Memorist Approach**

182 Imagine a person who has a remarkably retentive memory, whom we call a *Memorist*. When
 183 evaluating a program, the Memorist does all the computations eagerly; in the meantime, they
 184 are able to memorize, for each computation that would be thunked in a lazy evaluation, what
 185 other thunks (or more specifically, the evaluated results of which) it uses for its evaluation and the
 186 computation cost of this evaluation. At the end of evaluating the entire program, the Memorist learns
 187 all the thunks used by computing each part of the value, along with their individual computation
 188 cost. If someone or some program demands parts of the value, they can tell all the thunks that are
 189 actually used in the evaluation up to these parts. Now a corresponding lazy evaluation would also
 190 have to evaluate exactly these thunks, one can thus infer the lazy evaluation cost based on the
 191 information provided by the Memorist.
 192

193 *Thunks and annotations.* To realise this approach, we wrap every piece of data inside a *thunk*
 194 that tracks the data’s usage and cost during evaluations. Thunks here are intended to simulate
 195 thunks in lazy evaluation, but they are not encoded as suspended computations. Instead, each thunk

in the Memorist Semantics has a unique name to distinguish it from others, and an *annotation*, which is a pair of cost component and a set of thunk names. The former tracks the cost incurred by evaluating the thunk to its WHNF, *without* the cost contributed by evaluating any other thunks. The latter represents all “used thunks”, *i.e.*, thunks whose results are used in that evaluation. We refer to these sets of names of used thunks as *usage sets* in what follows. We annotate output values similarly, with the annotation corresponding to the cost and thunk usage incurred by evaluating the expression to its WHNF (given all needed thunks are evaluated).

For example, let’s consider the program `truePrefixOfAppend [true;false] [true]` from a Memorist’s perspective. To account for the demand of part of a list, we embed thunks into lists by putting both arguments to the constructor `cons` in thunks. We also thunk the entire lists when using them as function arguments, as function applications do not necessarily use their arguments. Representing a thunk with a name i and an annotation a by $x_i@a$, where x is the expression being thunked, the two input lists, after thunking, can be represented by

$$\left\{ \text{true}_{i_4@a_{i_4}} :: \left\{ \text{false}_{i_3@a_{i_3}} :: \left[\begin{smallmatrix} i_1 @ a_{i_1} \\ i_2 @ a_{i_2} \end{smallmatrix} \right] \right\}_{i_7@a_{i_7}} \right\}_{i_8@a_{i_8}} \quad \text{and} \quad \left\{ \text{true}_{i_6@a_{i_6}} :: \left[\begin{smallmatrix} i_5 @ a_{i_5} \\ i_8 @ a_{i_8} \end{smallmatrix} \right] \right\}_{i_8@a_{i_8}}$$

for some unique names i_1, \dots, i_8 and annotations a_{i_1}, \dots, a_{i_8} , with `::` denoting `cons`. We delay the formal definitions to [Section 3](#).

The Memorist Semantics in action. Under the Memorist Semantics, we evaluate the expression `truePrefixOfAppend [true;false] [true]` eagerly and track thunk usage and cost in the meantime. We show the detailed evaluation steps in [Fig. 3](#).

We first compute list appending ([Fig. 3a](#)). At each step, we need to “unthunk” the first argument of the function to access the list inside. For example, at the first step, we unthunk thunk i_7 wrapping the entire first argument. By doing so, we say that we *used* thunk i_7 at this step. Thunk i_7 is the only thunk we need to use at this step. We then recursively apply the function `append` to the tail of the first input list until we reach the empty list. When we reach the empty list, we need to unthunk both thunks i_1 and i_8 , and return the second input list. After this step, we wrap the returned result (bound to z_2 in the figure) in a new thunk with a fresh name j_1 .

The result thunked in j_1 is obtained from computing `append []_{i_1} { true_{i_6} :: []_{i_8} }`, which incurs a cost of 1 due to one call to `append` and uses the thunks i_1 and i_8 . Hence, its cost annotation is 1. For its usage set, we need to include the names i_1 and i_8 , as the corresponding thunks are used. Additionally, these two thunks also have their own thunk usage, recorded in the usage sets in their annotations a_{i_1} and a_{i_8} . We denote the usage sets annotated to these two thunks by s_{i_1} and s_{i_8} below respectively. These usage set annotations suggest that, if i_1 and i_8 are to be used, all the thunks in s_{i_1} and s_{i_8} must be used. To avoid repeated look up of usage set annotations subsequently, we propagate the usage of i_1 and i_8 by adding all theses thunk names to the usage set of j_1 . Below we denote the set of $\{i_1\} \cup s_{i_1}$ by $s(i_1)$ for short, and similar for other thunks. The usage set annotation to j_1 is then $s(i_1) \cup s(i_8)$.

The rest of the computation proceeds in a similar fashion. Note that when annotating each new thunk, we do not include the cost and thunk usage incurred by evaluating any thunks inside the values. For instance, the annotation to the thunk j_2 captures only the cost and thunk usage incurred *after* the thunk j_1 is created. By so doing, we effectively tracks the localised cost and usage to each thunk, so that the information recorded for a thunk k never includes the cost and usage from another thunk m unless the evaluated result of k is always needed to evaluate m . This will become helpful later when we analyse the lazy evaluation behaviour and cost. Lastly, the final output is also annotated in this way. We note that this final annotation represents exactly the thunk usage and cost due to evaluating `append [true;false] [true]` to its WHNF.

```

246      append  $\left\{ \text{true}_{i_4} :: \left\{ \text{false}_{i_3} :: \left[ \right]_{i_1 @ a_{i_1}} \right\}_{i_2 @ a_{i_2}} \right\}_{i_7 @ a_{i_7}} \left\{ \text{true}_{i_6} :: \left[ \right]_{i_5 @ a_{i_5}} \right\}_{i_8 @ a_{i_8}}$ 
247
248 = ✓let  $z_1 :=$  append  $\left\{ \text{false}_{i_3} :: \left[ \right]_{i_1} \right\}_{i_2} \left\{ \text{true}_{i_6} :: \left[ \right]_{i_5} \right\}_{i_8}$  in  $\text{true}_{i_4} :: \underline{z_1}$   $\rightsquigarrow i_7$ 
249
250 = let  $z_1 :=$  ✓ $($ let  $z_2 :=$  append  $\left[ \right]_{i_1} \left\{ \text{true}_{i_6} :: \left[ \right]_{i_5} \right\}_{i_8}$  in  $\text{false}_{i_3} :: \underline{z_2}$  $)$  in  $\text{true}_{i_4} :: \underline{z_1}$   $\rightsquigarrow i_2$ 
251 = let  $z_1 :=$  ✓ $($ let  $z_2 :=$  ✓ $\text{true}_{i_6} :: \left[ \right]_{i_5}$  in  $\text{false}_{i_3} :: \underline{z_2}$  $)$  in  $\text{true}_{i_4} :: \underline{z_1}$   $\rightsquigarrow i_1, i_8$ 
252 = let  $z_1 :=$   $\text{false}_{i_3} :: \left\{ \text{true}_{i_6} :: \left[ \right]_{i_5} \right\}_{j_1 @ (1, s(i_1) \cup s(i_8))}$  in  $\text{true}_{i_4} :: \underline{z_1}$ 
253
254 =  $\text{true}_{i_4} :: \left\{ \text{false}_{i_3} :: \left\{ \text{true}_{i_6} :: \left[ \right]_{i_5} \right\}_{j_1} \right\}_{j_2 @ (1, s(i_2))}$  @ $(1, s(i_7))$ 
255
256 (a) Evaluation steps of append
257    $\text{truePrefix} \left\{ \text{true}_{i_4} :: \left\{ \text{false}_{i_3} :: \left\{ \text{true}_{i_6} :: \left[ \right]_{i_5} \right\}_{j_1} \right\}_{j_2} \right\}_{j_3}$ 
258
259 = ✓let  $w_1 :=$   $\text{truePrefix} \left\{ \text{false}_{i_3} :: \left\{ \text{true}_{i_6} :: \left[ \right]_{i_5} \right\}_{j_1} \right\}_{j_2}$  in  $\text{true}_{i_4} \rightsquigarrow \text{true}_{i_4} :: \underline{w_1}$   $\rightsquigarrow j_3$ 
260
261 = let  $w_1 :=$  ✓ $($ let  $w_2 :=$   $\text{truePrefix} \left\{ \text{true}_{i_6} :: \left[ \right]_{i_5} \right\}_{j_1}$  in  $\text{false}_{i_3} \rightsquigarrow \text{false}_{i_3} :: \underline{w_2}$  $)$  in ...  $\rightsquigarrow j_2$ 
262 = let  $w_1 :=$  ✓ $($ let  $w_2 :=$  ✓ $($ let  $w_3 :=$   $\text{truePrefix} \left[ \right]_{i_5}$  in  $\text{true}_{i_6} \rightsquigarrow \text{true}_{i_6} :: \underline{w_3}$  $)$  in ... $)$  in ...  $\rightsquigarrow j_1$ 
263
264 = let  $w_1 :=$  ✓ $($ let  $w_2 :=$   $\left[ \right]$  in  $\text{true}_{i_6} \rightsquigarrow \text{true}_{i_6} :: \underline{w_3}$  $)$  in ...  $\rightsquigarrow i_5$ 
265
266 = let  $w_1 :=$  ✓ $($ let  $w_2 :=$   $\left( \text{true}_{i_6} \rightsquigarrow \text{true}_{i_6} :: \left[ \right]_{k_1 @ (1, s(i_5))} \right)$  in  $\text{false}_{i_3} \rightsquigarrow \text{false}_{i_3} :: \underline{w_2}$  $)$  in ...
267
268 = let  $w_1 :=$  ✓ $($ let  $w_2 :=$   $\text{true}_{i_6} :: \left[ \right]_{k_1}$  in  $\text{false}_{i_3} \rightsquigarrow \text{false}_{i_3} :: \underline{w_2}$  $)$  in  $\text{true}_{i_4} \rightsquigarrow \text{true}_{i_4} :: \underline{w_1}$   $\rightsquigarrow i_6$ 
269
270 where  $a_{k_2} = (1, s(j_1) \cup s(i_6)) = (1, \{j_1\} \cup s(i_1) \cup s(i_8) \cup s(i_6))$ 
271 = let  $w_1 := \left[ \right]$  in  $\text{true}_{i_4} \rightsquigarrow \text{true}_{i_4} :: \underline{w_1}$   $\rightsquigarrow i_3$ 
272 =  $\text{true}_{i_4} \rightsquigarrow \text{true}_{i_4} :: \left[ \right]_{k_3 @ (1, s(j_2) \cup s(i_3))}$ 
273 =  $\text{true}_{i_4} :: \left[ \right]_{k_3}$  @ $(1, s(j_3) \cup s(i_4))$   $\rightsquigarrow j_3, i_4$ 
274
275 (b) Evaluation steps of truePrefix
276    $\text{truePrefixOfAppend} \left\{ \text{true}_{i_4} :: \left\{ \text{false}_{i_3} :: \left[ \right]_{i_1} \right\}_{i_2} \right\}_{i_7} \left\{ \text{true}_{i_6} :: \left[ \right]_{i_5} \right\}_{i_8}$ 
277
278 = ✓let  $zs :=$  append  $\left\{ \text{true}_{i_4} :: \left\{ \text{false}_{i_3} :: \left[ \right]_{i_1} \right\}_{i_2} \right\}_{i_7} \left\{ \text{true}_{i_6} :: \left[ \right]_{i_5} \right\}_{i_8}$  in  $\text{truePrefix} \underline{zs}$ 
279
280 =  $\text{truePrefix} \left\{ \text{true}_{i_4} :: \left\{ \text{false}_{i_3} :: \left\{ \text{true}_{i_6} :: \left[ \right]_{i_5} \right\}_{j_1} \right\}_{j_2} \right\}_{j_3 @ (1, s(i_7))}$ 
281 =  $\text{true}_{i_4} :: \left[ \right]_{k_3}$  @ $(1 + 1, s(j_3) \cup s(i_4)) = (2, s(j_3) \cup s(i_4))$ 
282
283 (c) Evaluation steps of truePrefixOfAppend
284
285 Fig. 3. A stepwise illustration of the Memorist evaluation for the program truePrefixOfAppend. A tick (✓) is placed immediately after unfolding a step of function application to indicate that a cost is incurred. The notation  $w_1$  denotes the thunking of  $w_1$ , and  $e \rightsquigarrow xs$  is shorthand for if  $e$  then  $xs$  else  $[]$ . The notation  $\rightsquigarrow i, j, \dots$  denotes that thunks named  $i, j, \dots$  are used in the last step proceeding to the expression on the current line. Annotations to thunks already shown before are omitted in subsequent lines. The notation  $s(i)$  denotes the union of  $\{i\}$  and the usage set annotated to  $i$ . Annotations to the output are given at the final steps next to the output value.
286
287
288
289
290
291
292
293
294

```

295 *Lazy computation cost.* With the information in the output annotation and annotations, we can
 296 analyse the cost of append with respect to *any demand* on the output list. We can do so by collecting
 297 all the thunks used by a demand and aggregating their individual cost, as each thunk has already
 298 recorded the cost and thunk usage in its annotation.

299 For example, if only the first element in the output list of append is demanded, we only need
 300 to consider the thunks used to evaluate the output value to its outermost constructor, *i.e.*, $s(i_7)$
 301 in the annotation, the thunk i_4 wrapping the first element, and all the thunks used by i_4 . That is,
 302 we consider the set $s(i_4) \cup s(i_7)$. However, we are interested in the lazy cost of append itself, not
 303 in the previous computations leading to its input lists. Consequently, we do not want to account
 304 for the cost incurred by the evaluation to the values inside the thunks in the input. Instead, we
 305 remove all the thunks existing prior to the evaluation of append `[true; false] [true]`. In this case,
 306 we arrive at the empty set after such removal, suggesting there is no thunked cost we need to
 307 consider. Therefore, we only look at the cost annotation to the final output, which is 1 as shown in
 308 the last evaluation step of Fig. 3a. The inferred lazy cost is thus 1, as expected.

309 If the output is demanded to the first two elements instead, we need to additionally take the
 310 thunk j_2 wrapping the first cons cell and i_3 wrapping the second element in the output list into
 311 account. That is, we consider all thunks in the set $s(i_7) \cup s(i_4) \cup s(j_2) \cup s(i_3)$ with $s(j_2) = \{j_2\} \cup s(i_2)$.
 312 However, only j_2 in this set is created during the computation of append, which has a cost annotation
 313 of 1. Adding this to the output cost annotation gives us $1 + 1 = 2$, indicating that the lazy cost is 2
 314 under the demand in question. We can analyse cost for other demand in the same manner. Note
 315 that, while we have reasoned about different demand, we need not re-evaluate the entire expression,
 316 but only extract and aggregate information based on the demand from the same evaluated result.
 317

318 *Composing lazy cost analysis.* A key advantage of the Memorist Semantics is that we can analyze
 319 cost locally and compose cost analysis. To see why, let's consider the evaluation of `truePrefix`
 320 and `truePrefixOfAppend`. Figure 3b illustrates the evaluation of `truePrefix` applied to the result of
 321 append. As before, we wrap the input list in a thunk with a fresh name j_3 . The thunk j_3 also has its
 322 annotation which comes from some previous computations (for example, from thunking the result
 323 of append in computing `truePrefixOfAppend`). The evaluation of `truePrefix` proceeds in a similar
 324 manner as that of append.
 325

326 There are two things worth noticing in these evaluation steps. First, each time after returning
 327 from a recursive call, the value of the head element in the input list needs to be examined. If the
 328 head element is true, we cons the head element to the recursive result. Otherwise, we simply return
 329 the empty list. Accessing the value in the list requires *using* the thunk wrapping that list element.
 330 For example, consider the thunk k_2 wraps the result of computing $\text{true}_{i_6} \rightsquigarrow \text{true}_{i_6} :: []_{k_1}$ bound
 331 to w_2 . Computing the latter requires accessing the value inside the thunk i_6 to determine which
 332 branch to take for the if-expression. The thunk i_6 is used, and the usage $s(i_6)$ is thus taken as part
 333 of the usage set annotation to k_2 . A similar thing happens with k_3 and with the final output.

334 Second, if a part of a computation is eventually unnecessary, its cost and thunk usage will not
 335 be included anywhere in the final result. For example, thunks k_1 (wrapping the empty list) and k_2
 336 (wrapping the third element of the input list) are ‘thrown away’ after encountering `false`, so none
 337 of the associated cost or thunk usage is included in any part of the final output. Indeed, nothing
 338 after the second element in the input list is needed to compute the final result in lazy evaluation.

339 Now consider the top-level program `truePrefixOfAppend` in Fig. 3c. Evaluating the program
 340 `truePrefixOfAppend [true; false] [true]` under the Memorist Semantics essentially first computes
 341 append on the two lists (thunked and annotated), thunks the resulting list, and then computes
 342 `truePrefix` on that thunked list.

344 Types $A, B ::= \text{bool} \mid \text{list } A \mid A \times B \mid \text{T } A$
 345 Variables $x, y \in \text{Var}$
 346 Expressions $M, N ::= \text{true} \mid \text{false} \mid \text{if } M_1 \text{ then } M_2 \text{ else } M_3$
 347 $\mid (M, N) \mid \text{fst } M \mid \text{snd } M \mid x \mid \text{let } x = M \text{ in } N$
 348 $\mid \text{nil} \mid \text{cons } M N \mid \text{foldr } (\lambda xy. M_1) M_2 M_3$
 349 $\mid \text{tick } M \mid \text{lazy } M \mid \text{force } M$
 350

Fig. 4. Language Syntax

$$\begin{array}{c}
 \text{TBOOL} \quad \text{TNIL} \quad \text{TCONS} \\
 \frac{b \in \{\text{true}, \text{false}\}}{\gamma \vdash b : \text{bool}} \quad \frac{}{\gamma \vdash \text{nil} : \text{list } A} \quad \frac{\gamma \vdash M : \text{T } A \quad \gamma \vdash N : \text{T } (\text{list } A)}{\gamma \vdash \text{cons } M N : \text{list } A} \\
 \\
 \text{TPAIR} \quad \text{TFST} \quad \text{TSND} \quad \text{TVAR} \\
 \frac{\gamma \vdash M : A \quad \gamma \vdash N : B}{\gamma \vdash (M, N) : A \times B} \quad \frac{\gamma \vdash M : A \times B}{\gamma \vdash \text{fst } M : A} \quad \frac{\gamma \vdash M : A \times B}{\gamma \vdash \text{snd } M : B} \quad \frac{\gamma(x) = A}{\gamma \vdash x : A} \\
 \\
 \text{TIF} \quad \text{TLET} \\
 \frac{\gamma \vdash M_1 : \text{bool} \quad \gamma \vdash M_2 : A \quad \gamma \vdash M_3 : A}{\gamma \vdash \text{if } M_1 \text{ then } M_2 \text{ else } M_3 : A} \quad \frac{\gamma \vdash M : A \quad \gamma, x : A \vdash N : B}{\gamma \vdash \text{let } x = M \text{ in } N : B} \\
 \\
 \text{TFOLDR} \\
 \frac{\gamma, x : \text{T } A, y : \text{T } B \vdash M_1 : B \quad \gamma \vdash M_2 : B \quad \gamma \vdash M_3 : \text{list } A}{\gamma \vdash \text{foldr } (\lambda xy. M_1) M_2 M_3 : B} \\
 \\
 \text{TTICK} \quad \text{TLAZY} \quad \text{TFORCE} \\
 \frac{\gamma \vdash M : A}{\gamma \vdash \text{tick } M : A} \quad \frac{\gamma \vdash M : A}{\gamma \vdash \text{lazy } M : \text{T } A} \quad \frac{\gamma \vdash M : \text{T } A}{\gamma \vdash \text{force } M : A}
 \end{array}$$

Fig. 5. Typing Rules

375 We can then perform thunk collection based on demand as before. If the output list is demanded
 376 only to its first element, the set of thunks collected is $s(j_3) \cup s(i_4)$, among which only j_3 is not
 377 presented in the input environment of this program. The cost is thus the sum of the cost annotation
 378 given to the final output and to the thunk j_3 , which is $2 + 1 = 3$. If the demand is to construct the
 379 entire list, we take also $s(k_3)$ into account, where k_3 is the thunk wrapping the empty tail. With
 380 k_3 and j_2 being the two thunks not presented in the input alongside j_3 and each having a cost
 381 annotation of 1, the cost is inferred to be $2 + 1 + 1 + 1 = 5$.
 382

383 Notice that, regardless of the demand, the cost and thunk usage associated with computing
 384 the base case of append will not be considered in this instance, aligning with the fact that a lazy
 385 evaluation of the program will not process the rest of the first input list after the element false.
 386

3 The Memorist Semantics

We now formally define the Memorist Semantics.

393 3.1 Language Syntax and Type System

394 We define the Memorist semantics for a typed total language \mathcal{L} with booleans, lists, pairs, explicit
 395 thunks, ticks, and structural recursion on lists, similar to the language considered by Xia et al. [2024].
 396 Its syntax and typing rules of the language are defined in Fig. 4 and 5, where well-typed terms are
 397 given by judgements of the form $\gamma \vdash M : A$. Thunks are represented explicitly with the T type. The
 398 constructs lazy and force manipulates thunks explicitly, with the former thunking a computation,
 399 and the latter forcing the evaluation of the thunk. The tick construct simulates a computation
 400 that incurs a unit cost. The construct foldr is included as a primitive of the language to facilitate
 401 structural recursions on lists. The language can be regarded as an intermediate representation into
 402 which an ordinary functional language can be translated. Having thunks as a syntactic construct
 403 enables representing and studying laziness directly.

404 3.2 Memorist Semantics

405 We define the Memorist semantics for the language \mathcal{L} formally below. Under the Memorist Semantics,
 406 terms are evaluated into values, which are defined by

$$407 \text{Value } v ::= \text{true} \mid \text{false} \mid \text{nil} \mid \text{cons } v_1 v_2 \mid (v_1, v_2) \mid \text{thunk}_i v$$

408 with an additional typing rule

$$\frac{\gamma \vdash v : A}{\gamma \vdash \text{thunk}_i v : \text{T}\ A} \quad \text{TMTHUNK}$$

409 Thunks are represented by values of the thunk type. As previously described in Section 2, each
 410 thunk is assigned a unique name $i \in \mathbb{N}$ when created to distinguish it from other thunks.

411 The semantics associates the output value and each thunk with an annotation that records cost
 412 and thunk usage related information. Each annotation α is a pair (c, s) of a cost annotation c and a
 413 set s containing the names of thunks. The cost annotation c is the cost accumulated during the
 414 evaluation that is never thunked by the end of the evaluation. As already can be seen from the
 415 motivating example presented earlier, the cost annotation c is typically only part of the cost of
 416 evaluating an expression to a value. To derive the evaluation cost, we must at least consider c
 417 together with thunks whose names are in the set s . The set s in the annotation is a collection of
 418 names of all the thunks whose evaluation must be forced to arrive at the WHNF of the value in
 419 question. We call this set of a *usage set* of this value. Each component of an annotation can be
 420 extracted via the usual first and second projections on pairs, denoted by $\pi_1(\cdot)$ and $\pi_2(\cdot)$ respectively.

421 An evaluation occurs in an evaluation environment, defined as $\Gamma ::= \emptyset \mid \Gamma, x \mapsto v$, that maps
 422 variables to values. We also define a *annotation context*, $\mathcal{A} ::= \emptyset \mid \mathcal{A}, i \mapsto \alpha_i$, that maps a name i
 423 of a thunk to its annotation α_i , in order to track and record the annotations so that they can be
 424 referred to later. In other words, two separate contexts are kept track of in the evaluation. This
 425 allows more flexibility in manipulating thunks and analysing their usage and cost.

426 Formally, the Memorist semantics for well-typed terms in \mathcal{L} is defined by the evaluation judgement
 427 $\Gamma; \mathcal{A} \vdash M \Downarrow \mathcal{A}' ; \langle v, \alpha \rangle$. It states that the well-typed term M evaluates, under the environment
 428 Γ and the annotation context \mathcal{A} , to the value v annotated by α with the extended annotation
 429 context \mathcal{A}' . This judgement is defined by the operational semantic rules presented in Figure 6.
 430 The evaluation itself is performed eagerly by default, which is easier to reason about than its lazy
 431 counterpart. We claim (and prove in the next section) that the recorded thunk usage reflects lazy
 432 behaviours, in the sense that it captures exactly the thunks needed to be evaluated and no more in
 433 a corresponding lazy evaluation, from which we can derive the lazy evaluation cost.

434 The EMBASIC rule introduces straightforwardly the base cases. The EMVAR rule looks up the
 435 variable in the environment for the value it binds to. The rules EMCONS, EMPAIR and EMLET

436

442 EMBASIC 443 $t \in \{\text{true}, \text{false}, \text{nil}_A\}$ 444 $\frac{}{\Gamma; \mathcal{A} \vdash t \Downarrow \mathcal{A}; \langle t, \odot \rangle}$	444 EMCONS 445 $\frac{\Gamma; \mathcal{A} \vdash M \Downarrow \mathcal{A}_1; \langle v_1, \alpha_1 \rangle \quad \Gamma; \mathcal{A}_1 \vdash N \Downarrow \mathcal{A}_2; \langle v_2, \alpha_2 \rangle}{\Gamma; \mathcal{A} \vdash \text{cons } M N \Downarrow \mathcal{A}_2; \langle \text{cons } v_1 v_2, \alpha_1 \oplus \alpha_2 \rangle}$
446 EMIFTRUE 447 $\frac{\Gamma; \mathcal{A} \vdash M_1 \Downarrow \mathcal{A}_1; \langle \text{true}, \alpha_1 \rangle \quad \Gamma; \mathcal{A}_1 \vdash M_2 \Downarrow \mathcal{A}_2; \langle v, \alpha_2 \rangle}{\Gamma; \mathcal{A} \vdash \text{if } M_1 \text{ then } M_2 \text{ else } M_3 \Downarrow \mathcal{A}_2; \langle v, \alpha_1 \oplus \alpha_2 \rangle}$	446 EMFST 447 $\frac{\Gamma; \mathcal{A} \vdash M \Downarrow \mathcal{A}_1; \langle (v_1, v_2), \alpha_1 \rangle}{\Gamma; \mathcal{A} \vdash \text{fst } M \Downarrow \mathcal{A}_1; \langle v_1, \alpha_1 \rangle}$
448 EMIFFFALSE 449 $\frac{\Gamma; \mathcal{A} \vdash M_1 \Downarrow \mathcal{A}_1; \langle \text{false}, \alpha_1 \rangle \quad \Gamma; \mathcal{A}_1 \vdash M_3 \Downarrow \mathcal{A}_2; \langle v, \alpha_2 \rangle}{\Gamma; \mathcal{A} \vdash \text{if } M_1 \text{ then } M_2 \text{ else } M_3 \Downarrow \mathcal{A}_2; \langle v, \alpha_1 \oplus \alpha_2 \rangle}$	448 EMSND 449 $\frac{\Gamma; \mathcal{A} \vdash M \Downarrow \mathcal{A}_1; \langle (v_1, v_2), \alpha_1 \rangle}{\Gamma; \mathcal{A} \vdash \text{snd } M \Downarrow \mathcal{A}_1; \langle v_2, \alpha_1 \rangle}$
450 EMPAIR 451 $\frac{\Gamma; \mathcal{A} \vdash M \Downarrow \mathcal{A}_1; \langle v_1, \alpha_1 \rangle \quad \Gamma; \mathcal{A}_1 \vdash N \Downarrow \mathcal{A}_2; \langle v_2, \alpha_2 \rangle}{\Gamma; \mathcal{A} \vdash (M, N) \Downarrow \mathcal{A}_2; \langle (v_1, v_2), \alpha_1 \oplus \alpha_2 \rangle}$	450 EMTICK 451 $\frac{}{\Gamma; \mathcal{A} \vdash \text{tick } M \Downarrow \mathcal{A}'; \langle v, \alpha \rangle}$
452 EMVAR 453 $\frac{\Gamma(x) = v}{\Gamma; \mathcal{A} \vdash x \Downarrow \mathcal{A}; \langle v, \odot \rangle}$	452 EMLET 453 $\frac{\Gamma; \mathcal{A} \vdash M \Downarrow \mathcal{A}_1; \langle v_1, \alpha_1 \rangle \quad \Gamma, (x \mapsto v_1); \mathcal{A}_1 \vdash N \Downarrow \mathcal{A}_2; \langle v_2, \alpha_2 \rangle}{\Gamma; \mathcal{A} \vdash \text{let } x = M \text{ in } N \Downarrow \mathcal{A}_2; \langle v_2, \alpha_1 \oplus \alpha_2 \rangle}$
454 EMFOLDRNIL 455 $\frac{\Gamma; \mathcal{A} \vdash M_3 \Downarrow \mathcal{A}_1; \langle \text{nil}_A, \alpha_1 \rangle \quad \Gamma; \mathcal{A}_1 \vdash M_2 \Downarrow \mathcal{A}_2; \langle v, \alpha_2 \rangle}{\Gamma; \mathcal{A} \vdash \text{foldr } (\lambda xy. M_1) M_2 M_3 \Downarrow \mathcal{A}_2; \langle v, \alpha_1 \oplus \alpha_2 \rangle}$	
456 EMFOLDRCONS 457 $\frac{\Gamma; \mathcal{A} \vdash M_3 \Downarrow \mathcal{A}_1; \langle \text{cons } v (\text{thunk}_i vs), \alpha_1 \rangle \quad \Gamma; \mathcal{A}_1 \vdash \text{foldr } (\lambda xy. M_1) M_2 vs \Downarrow \mathcal{A}_2; \langle v_2, \alpha_2 \rangle}{\begin{aligned} j \notin \text{dom}(\mathcal{A}_2) \cup \{i\} & \quad x' \neq y' \quad x', y' \notin \text{dom}(\Gamma) \cup \text{FV}(M_1) \cup \text{FV}(M_2) \cup \text{FV}(M_3) \\ \Gamma, (x' \mapsto v), (y' \mapsto \text{thunk}_j v_2); \mathcal{A}_2, (j \mapsto \alpha_2 \oplus \{i\} \oplus \pi_2(\mathcal{A}_2(i))) \vdash M_1[x', y'/x, y] \Downarrow \mathcal{A}_3; \langle v_3, \alpha_3 \rangle \end{aligned}}$	456 EMFORCE 457 $\frac{\Gamma; \mathcal{A} \vdash M \Downarrow \mathcal{A}'; \langle v, \alpha \rangle \quad i \notin \text{dom}(\mathcal{A}')}{\Gamma; \mathcal{A} \vdash \text{force } M \Downarrow \mathcal{A}'; \langle v, \alpha \oplus \{i\} \oplus \pi_2(\mathcal{A}'(i)) \rangle}$
458 EMLAZY 459 $\frac{\Gamma; \mathcal{A} \vdash M \Downarrow \mathcal{A}'; \langle v, \alpha \rangle \quad i \notin \text{dom}(\mathcal{A}')}{\Gamma; \mathcal{A} \vdash \text{lazy } M \Downarrow \mathcal{A}', (i \mapsto \alpha); (\text{thunk}_i v, \odot)}$	

Fig. 6. The Memorist Semantics. The notation $\text{FV}(M)$ denotes the set of all free variables appearing in M . The empty annotation $(0, 0)$ is denoted \odot . The operation \oplus on annotations evaluates to an annotation composed of the sum of the costs and the union of the sets of the two input annotations. That is, $\alpha_1 \oplus \alpha_2 := (\pi_1(\alpha_1) + \pi_1(\alpha_2), \pi_2(\alpha_1) \cup \pi_2(\alpha_2))$. We use the shorthand $\alpha \oplus n$ where n is a number to mean $(\pi_1(\alpha) + n, \pi_2(\alpha))$ and $\alpha \oplus s$ where s is a set to mean $(\pi_1(\alpha), \pi_2(\alpha) \cup s)$.

proceed by evaluating the subterms separately but consecutively. The annotations to values of the two sub-evaluations are combined by adding the cost and taking the union of the resulting usage sets. With set union, the semantics is able to account for potential sharing of thunks. The rules EMIFTRUE and EMIFFFALSE follow the similar vein, though the then- and else-branches are evaluated only when the if-condition evaluates to true and false respectively. The rules for pair projections, EMFST and EMSND, evaluate a term M to a value of pair and return the corresponding component. The rule EMTICK increments the cost count by one.

491 In the EMLAZY rule, the term M is evaluated to a value before being wrapped inside a thunk
 492 constructor. The computation cost and usage set annotation associated with this evaluation is
 493 annotated to the thunked value v that M evaluates to, by extending the annotation context with this
 494 annotation for the introduced thunk. Accordingly, the cost and usage set associated with evaluating
 495 the halted computation thunk $_i v$ is empty, as a thunked computation incurs no cost and utilises
 496 no additional thunks per se. Its associated cost and usage set are taken into effect only when the
 497 thunk is needed, forcing the computation to actually take place.

498 Forcing of the evaluation of thunks is explicitly done via force in this language, which is evaluated
 499 according to the EMFORCE rule. While thunks in the usage set annotated to the thunked value
 500 (i.e., $\pi_2(\mathcal{A}'(i))$ in the rule) are directly merged into the usage set annotated to the final value via
 501 set union, the thunked cost is not added in yet at this point. This avoids the duplication of cost
 502 when the same thunk is needed more than once in the evaluation. The total evaluation cost is to be
 503 derived afterwards, based on the output usage set and the annotations.

504 The term foldr $(\lambda x y . M_1) M_2 M_3$ is evaluated based on whether or not M_3 is the empty list. When
 505 the argument list is empty, the result is simply the value of M_2 , as specified by the rule EMFOLDRNIL.
 506 Otherwise, the evaluation follows the rule EMFOLDRCONS, which first recursively evaluates foldr
 507 on the tail of the list, and then applies $\lambda x y . M_1$ to the head of the list and the result of the recursive
 508 call. Notice that the results of the recursive calls are thunked: the step function $\lambda x y . M_1$ does not
 509 always have to use any of its argument.

510 We are only interested in annotation contexts that are *valid* in the following sense. Below let
 511 $\text{dom}(\cdot)$ and $\text{im}(\cdot)$ denote the domain and the image of a function, respectively.

512 *Definition 3.1 (Valid Memorist annotation contexts).* An annotation context \mathcal{A} is *valid* if for every
 513 annotation $(c, s) \in \text{im}(\mathcal{A})$, we have $s \subseteq \text{dom}(\mathcal{A})$. An annotation context \mathcal{A} is *valid* for a value v if
 514 it is itself valid, and for any name i assigned to a thunk inside v , $i \in \text{dom}(\mathcal{A})$. Moreover, \mathcal{A} is *valid*
 515 for an environment Γ if $\forall x \in \text{dom}(\Gamma)$, \mathcal{A} is valid for $\Gamma(x)$.

516 By ‘thunks inside v ’, we mean thunks wrapping the arguments to the (outermost and nested)
 517 constructors of v . These include, for instance, all four thunks appearing in the list

$$518 \quad \text{cons} (\text{thunk}_i \text{ true}) (\text{thunk}_j (\text{cons} (\text{thunk}_k \text{ false}) (\text{thunk}_m \text{ nil}))).$$

519 *Basic Properties of the Annotation Context.* The Memorist semantics has the following properties.
 520 An evaluation may extend the annotation context, but it never modifies the existing annotations.
 521 Furthermore, validity of annotation contexts is preserved by evaluation.

522 *LEMMA 3.2 (EVALUATION ONLY EXTENDS ANNOTATION CONTEXT).* If $\Gamma; \mathcal{A} \vdash M \Downarrow \mathcal{A}'; \langle v, \alpha \rangle$ then
 523 $\forall i \in \text{dom}(\mathcal{A}), \mathcal{A}'(i) = \mathcal{A}(i)$.

524 *LEMMA 3.3 (EVALUATION PRESERVES ANNOTATION CONTEXT VALIDITY).* If $\Gamma; \mathcal{A} \vdash M \Downarrow \mathcal{A}'; \langle v, \alpha \rangle$ and \mathcal{A} is valid for Γ , we have that \mathcal{A}' is valid for Γ and v , and that $\pi_2(\alpha) \subseteq \text{dom}(\mathcal{A}')$.

525

3.3 Lazy Cost and Usage Analysis using Annotations Derived from Memorist Evaluation

 526 In the Memorist semantics, each annotation is uniquely associated with the outermost portion
 527 of the annotated value. Hence, to derive the complete usage set and cost due to evaluating the
 528 expression to this value, we must consider every thunk that is inferred to be needed. Consider an
 529 evaluation $\Gamma; \mathcal{A} \vdash M \Downarrow \mathcal{A}'; \langle v, (c, s) \rangle$. As (c, s) is the annotation given to the value v , all the thunks
 530 recorded in the usage set s are required in order to lazily evaluate to the outermost portion of v .
 531 The total lazy cost of the evaluation, i.e., the cost that would be incurred if the evaluation were
 532 call-by-need, is then the sum of all annotated costs given to every thunk in s , together with c .

540 To perform such calculations, we define the following operation

$$541 \quad \text{sumcost}_M(\mathcal{A}, s) := \sum_{i \in s \cap \text{dom}(\mathcal{A})} \pi_1(\mathcal{A}(i))$$

542 for an annotation context \mathcal{A} and a usage set s . The total lazy cost of the above evaluation to the
 543 outermost portion of v (i.e., to the expression in WHNF) can be then expressed as

$$544 \quad c + \text{sumcost}_M(\mathcal{A}' \setminus \mathcal{A}, s).$$

545 We use set difference on the extended and initial annotation contexts (\mathcal{A}' and \mathcal{A} respectively,
 546 treated as functions, i.e, set of pairs of mappings) to exclude thunks already evaluated prior to the
 547 evaluation we are considering. Set difference is essential; otherwise, we would also add in any cost
 548 incurred by computations prior to the current one.

549 Sometimes more than the outermost portion of a value may be used. For example, suppose the
 550 above evaluation yields

$$551 \quad \langle \text{cons} (\text{thunk}_i \text{ true}) (\text{thunk}_j \text{ nil}), \quad (c, \{k, m\}) \rangle,$$

552 and that it occurs in a larger context where the value of the head of the list is eventually also used.
 553 In this case, the needed thunks are not simply the thunks k and m , but also the thunk i and any
 554 further thunks used by the evaluation to the value inside that thunk. In other words, the complete
 555 usage set with respect to the demand in question should be $s' = \{k\} \cup \{i\} \cup \pi_2((\mathcal{A}'(i)))$. The cost
 556 can then be computed accordingly by $c + \text{sumcost}_M(\mathcal{A}' \setminus \mathcal{A}, s')$.

557 When the set s' reflects exactly what is needed to evaluate under a call-by-need strategy, the cost
 558 computed as above will result in the same cost as that produced by the corresponding call-by-need
 559 evaluation. In particular, collecting all the relevant names into a set before adding up the cost
 560 ensures that no thunk can contribute to the total cost more than once. This avoids duplicating the
 561 cost of a thunk even if it is used multiple times. It also explains why we collect the name of a forced
 562 thunk (along with names in its annotated usage set) into the usage set in the annotation given to
 563 the final value without adding in the thunked cost immediately in the evaluation rule EMFORCE:
 564 we will not duplicate the cost if the thunk is forced more than once in the evaluation.

572 4 Correctness of the Memorist Semantics

573 We prove that the Memorist Semantics is correct, by relating it to the Clairvoyance Semantics.
 574 Initially introduced by Hackett and Hutton [2019], the Clairvoyance Semantics models call-by-need
 575 evaluation with nondeterminism, and is shown in the same paper to be equivalent to the standard
 576 call-by-need semantics [Launchbury 1993]. The Clairvoyance Semantics nondeterministically
 577 chooses to proceed with or skip the evaluation of an expression when it is first encountered, instead
 578 of delaying the evaluation altogether until it is actually needed (if at all).

579 Reasoning about the relation between the Memorist semantics and the Clairvoyance Semantics
 580 directly is challenging. Instead, we define, as a stepping stone, a variant of the Clairvoyance
 581 Semantics in which thunks are named and have their own cost annotations. We prove that the
 582 annotated variant is equivalent to the original one, and that the Memorist semantics corresponds
 583 to this annotated variant. This correspondence takes evaluation cost into account; hence it shows
 584 the Memorist Semantics not only evaluates correctly but also infers correctly the lazy cost. All
 585 theorems and proofs presented here have been formalized in Rocq Prover, the code of which can
 586 be found in the accompanying artefact.

589 4.1 The Clairvoyance Semantics

590 Here we present a Clairvoyance Semantics to the language \mathcal{L} , adapted directly from the formalisation
 591 by Xia et al. [2024] which is itself based on a monadic variant of the Clairvoyance Semantics
 592 due to [Li et al. 2021]. Types are interpreted as

$$\begin{aligned} 593 \llbracket A \rrbracket &: \text{Set} \\ 594 \llbracket \text{bool} \rrbracket &:= \{\text{true}, \text{false}\} \\ 595 \llbracket \text{list } A \rrbracket &:= \{\text{nil}\} \cup \{\text{cons } \hat{v}_1 \hat{v}_2 \mid \hat{v}_1 \in \llbracket T A \rrbracket, \hat{v}_2 \in \llbracket T (\text{list } A) \rrbracket\} \\ 596 \llbracket A \times B \rrbracket &:= \{(\hat{v}_1, \hat{v}_2) \mid \hat{v}_1 \in \llbracket A \rrbracket, \hat{v}_2 \in \llbracket B \rrbracket\} \\ 597 \llbracket T A \rrbracket &:= \{\perp\} \cup \{\text{thunk } \hat{v} \mid \hat{v} \in \llbracket A \rrbracket\} \end{aligned}$$

600 The thunk type T is interpreted as a set whose elements are either of the form $\text{thunk } \hat{v}$ representing
 601 a thunk evaluated to a value v , or \perp representing skipped computation. The interpretation extends
 602 to the type context Γ , so that an element of the interpreted $\llbracket \Gamma \rrbracket$ is an environment of the evaluation.
 603

604 The semantics of evaluating a well-typed term $\gamma \vdash M : A$ is denoted by

$$605 \llbracket M \rrbracket : \llbracket \gamma \rrbracket \rightarrow \mathcal{P}(\llbracket A \rrbracket \times \mathbb{N})$$

606 that takes an interpreted environment $\hat{\Gamma} \in \llbracket \gamma \rrbracket$ and produces a set of pairs of values $\hat{v} \in \llbracket A \rrbracket$ and
 607 numbers $c \in \mathbb{N}$ of the cost incurred by the evaluation. The detailed definition of the monadic
 608 semantics can be found in Li et al. [2021]; Xia et al. [2024]. The output is denoted as a set, as the
 609 evaluation is nondeterministic. When evaluating lazy M and foldr which involve creating thunks,
 610 the semantics nondeterministically chooses to evaluate a subterm to a value and wrap it in the thunk
 611 constructor to represent an evaluated thunk, or skips the evaluation by producing the placeholder
 612 \perp . As the nondeterministic choice is done immediately when a thunk is created, forcing a thunk
 613 via the term force N is then simply accessing the already evaluated value in the thunk, if thunk is
 614 evaluated at all when created. The evaluation of force N fails if the computation of N is skipped.
 615

616 In addition to computing the values, the semantics also keeps track of cost, which is the number
 617 c in the output of the function $\llbracket \cdot \rrbracket$ on terms. This is the total evaluation cost, which varies along
 618 with the nondeterministic choices made during the evaluation. For a given term and environment,
 619 if the output is nonempty, the minimal cost for the same value corresponds to the call-by-need cost.

620 4.2 A Cost Annotated Variant of the Clairvoyance Semantics

621 We define a new variant of the Clairvoyance Semantics that is equivalent to the monadic Clairvoyance
 622 Semantics, and then establish the correspondence of the Memorist Semantics with the new
 623 variant. In this variant, we name thunks, and give and the final value cost annotations, similar to
 624 the Memorist Semantics except there are no usage sets. Values are defined as
 625

$$626 \tilde{v} ::= \text{true} \mid \text{false} \mid \text{nil} \mid \text{cons } \tilde{v}_1 \tilde{v}_2 \mid (\tilde{v}_1, \tilde{v}_2) \mid \text{thunk}_i \tilde{v} \mid \perp$$

627 with additional typing rules

$$\frac{\gamma \vdash \tilde{v} : A}{\gamma \vdash \text{thunk}_i \tilde{v} : T A} \quad \text{TATHUNK} \qquad \frac{}{\gamma \vdash \perp : T A} \quad \text{TABOT}$$

632 where $i \in \mathbb{N}$ is a unique name assigned to a thunk. Similar to the monadic Clairvoyance Semantics
 633 but unlike the Memorist Semantics, values of the thunk type T can take two forms: either an
 634 evaluated thunk, $\text{thunk}_i v$, or a skipped computation \perp .

635 Evaluations occur in an environment defined as $\tilde{\Gamma} ::= \emptyset \mid \tilde{\Gamma}, x \mapsto \tilde{v}$. We also define a cost
 636 annotation context $C ::= \emptyset \mid C, i \mapsto c_i$ that maps thunk names to the corresponding annotations.
 637

<p>638 EABASIC 639 $t \in \{\text{true}, \text{false}, \text{nil}\}$</p>	<p>640 EACONS 641 $\frac{\tilde{\Gamma}; C \vdash M \Downarrow^A C_1; \langle \tilde{v}_1, c_1 \rangle \quad \tilde{\Gamma}; C_1 \vdash N \Downarrow^A C_2; \langle \tilde{v}_2, c_2 \rangle}{\tilde{\Gamma}; C \vdash \text{cons } M N \Downarrow^A C_2; \langle \text{cons } \tilde{v}_1 \tilde{v}_2, c_1 + c_2 \rangle}$</p>
<p>642 EAIfTRUE 643 $\tilde{\Gamma}; C \vdash M_1 \Downarrow^A C_1; \langle \text{true}, c_1 \rangle \quad \tilde{\Gamma}; C_1 \vdash M_2 \Downarrow^A C_2; \langle \tilde{v}, c_2 \rangle$</p>	<p>644 EAfst 645 $\frac{\tilde{\Gamma}; C \vdash M \Downarrow^A C_1; \langle (\tilde{v}_1, \tilde{v}_2), c \rangle}{\tilde{\Gamma}; C \vdash \text{fst } M \Downarrow^A C_1; \langle \tilde{v}_1, c \rangle}$</p>
<p>646 EAIfFALSE 647 $\tilde{\Gamma}; C \vdash M_1 \Downarrow^A C_1; \langle \text{false}, c_1 \rangle \quad \tilde{\Gamma}; C_1 \vdash M_3 \Downarrow^A C_3; \langle \tilde{v}, c_2 \rangle$</p>	<p>648 EASND 649 $\frac{\tilde{\Gamma}; C \vdash M \Downarrow^A C_1; \langle (\tilde{v}_1, \tilde{v}_2), c \rangle}{\tilde{\Gamma}; C \vdash \text{snd } M \Downarrow^A C_1; \langle \tilde{v}_2, c \rangle}$</p>
<p>650 EAFORCE 651 $\tilde{\Gamma}; C \vdash M \Downarrow^A C'; \langle \text{thunk}_i \tilde{v}, c \rangle$</p>	<p>652 EAPAIR 653 $\frac{\tilde{\Gamma}; C \vdash M \Downarrow^A C_1; \langle \tilde{v}_1, c_1 \rangle \quad \tilde{\Gamma}; C_1 \vdash N \Downarrow^A C_2; \langle \tilde{v}_2, c_2 \rangle}{\tilde{\Gamma}; C \vdash (M, N) \Downarrow^A \mathcal{A}_2; \langle (\tilde{v}_1, \tilde{v}_2), c_1 + c_2 \rangle}$</p>
<p>654 EALAZY 655 $\frac{\tilde{\Gamma}; C \vdash M \Downarrow^A C'; \langle \tilde{v}, c \rangle \quad i \notin \text{dom}(C')}{\tilde{\Gamma}; C \vdash \text{lazy } M \Downarrow^A C', (i \mapsto c); \langle \text{thunk}_i \tilde{v}, 0 \rangle}$</p>	<p>656 EALAZYSKIP 657 $\frac{}{\tilde{\Gamma}; C \vdash \text{lazy } M \Downarrow^A C; \langle \perp, 0 \rangle}$</p>
<p>658 EAVAR 659 $\tilde{\Gamma}(x) = \tilde{v}$</p>	<p>660 EALET 661 $\frac{\tilde{\Gamma}; C \vdash M \Downarrow^A C_1; \langle \tilde{v}_1, c_1 \rangle \quad (\tilde{\Gamma}, x \mapsto \tilde{v}_1); C_1 \vdash N \Downarrow^A C_2; \langle \tilde{v}_2, c_2 \rangle}{\tilde{\Gamma}; C \vdash \text{let } x = M \text{ in } N \Downarrow^A C_2; \langle \tilde{v}_2, c_1 + c_2 \rangle}$</p>
<p>662 EATICK 663 $\tilde{\Gamma}; C \vdash M \Downarrow^A C'; \langle \tilde{v}, c \rangle$</p>	<p>664 EAFOLDRNIL 665 $\frac{\tilde{\Gamma}; C \vdash M_3 \Downarrow^A C_1; \langle \text{nil}, c_1 \rangle \quad \tilde{\Gamma}; C_1 \vdash M_2 \Downarrow^A C_2; \langle \tilde{v}, c_2 \rangle}{\tilde{\Gamma}; C \vdash \text{foldr } (\lambda x y. M_1) M_2 M_3 \Downarrow^A C_2; \langle \tilde{v}, c_1 + c_2 \rangle}$</p>
<p>666 EAFOLDRCONS 667 $\tilde{\Gamma}; C \vdash M_3 \Downarrow^A C_1; \langle \text{cons } \tilde{v} (\text{thunk}_i \tilde{v}s), c_1 \rangle \quad \tilde{\Gamma}; C_1 \vdash \text{foldr } (\lambda x y. M_1) M_2 \tilde{v}s \Downarrow^A C_2; \langle \tilde{v}_2, c_2 \rangle$</p>	<p>668 $j \notin \text{dom}(C_2) \cup \{i\} \quad x' \neq y' \quad x', y' \notin \text{dom}(\tilde{\Gamma}) \cup \text{FV}(M_1) \cup \text{FV}(M_2) \cup \text{FV}(M_3)$</p>
<p>669 $(\tilde{\Gamma}, (x' \mapsto \tilde{v}), (y' \mapsto \text{thunk}_i \tilde{v}_2)); C_2, (j \mapsto c_2) \vdash M_1[x', y'/x, y] \Downarrow^A C_3; \langle \tilde{v}_3, c_3 \rangle$</p>	<p>670 $\frac{}{\tilde{\Gamma}; C \vdash \text{foldr } (\lambda x y. M_1) M_2 M_3 \Downarrow^A C_3; \langle \tilde{v}_3, c_1 + c_3 \rangle}$</p>
<p>671 EAFOLDRCONSSKIP 672 $\tilde{\Gamma}; C \vdash M_3 \Downarrow^A C_1; \langle \text{cons } \tilde{v} \tilde{v}s, c_1 \rangle \quad x' \neq y' \quad x', y' \notin \text{dom}(\tilde{\Gamma}) \cup \text{FV}(M_1) \cup \text{FV}(M_2) \cup \text{FV}(M_3)$</p>	<p>673 $\frac{\tilde{\Gamma}, (x' \mapsto \tilde{v}_1), (y' \mapsto \perp); C_2 \vdash M_1[x', y'/x, y] \Downarrow^A C_3; \langle \tilde{v}_2, c_3 \rangle}{\tilde{\Gamma}; C \vdash \text{foldr } (\lambda x y. M_1) M_2 M_3 \Downarrow^A C_3; \langle \tilde{v}_2, c_1 + c_3 \rangle}$</p>

Fig. 7. The Annotated Clairvoyance Semantics. The set of free variables in M is denoted $\text{FV}(M)$.

The semantics is defined by evaluation judgements of the form $\tilde{\Gamma}; C \vdash M \Downarrow^A C'; \langle \tilde{v}, c \rangle$. It states that a well-typed term M is evaluated in an environment $\tilde{\Gamma}$ and a cost annotation context C to the value \tilde{v} with a cost annotation c and with the cost annotation context C' . The cost annotations here are essentially the same thing as in the Memorist Semantics.

We show all the semantics rules in Fig. 7. Most rules are similar to the evaluation rules of the Memorist Semantics, sans the tracking of thunk usage with sets. The main difference resides in the evaluation of lazy and foldr, where thunks are created. Under Clairvoyance Semantics, the evaluation nondeterministically evaluates the term and wraps the result with a thunk (rules EALAZY and EAFLDRCONS), or skips the evaluation and outputs \perp (rule EALAZYSKIP and EAFLDRCONSSKIP). Likewise, the evaluation of a term force M only succeeds if M evaluates to an evaluated thunk, and fails if the evaluation gives \perp that represents a skipped computation.

Basic properties of the cost annotation context. An evaluation under the Annotated Clairvoyance Semantics may extend the cost annotation context but never modifies the existing annotations.

LEMMA 4.1 (EVALUATION ONLY EXTENDS COST ANNOTATION CONTEXT). *If $\tilde{\Gamma}; C \vdash M \Downarrow^A C'; \langle \tilde{v}, c \rangle$, then $\forall i \in \text{dom}(C), C'(i) = C(i)$.*

We can analogously define a notion of validity of the Clairvoyance cost annotation contexts and show that it is preserved by evaluation.

Definition 4.2 (Valid Clairvoyance cost annotation context). An cost annotation context C is *valid* for a value \tilde{v} , if for any name i assigned to a thunk nested inside \tilde{v} , $i \in \text{dom}(C)$. Moreover, C is *valid* for an environment $\tilde{\Gamma}$ if $\forall x \in \text{dom}(\tilde{\Gamma}), C$ is valid for $\tilde{\Gamma}(x)$.

LEMMA 4.3 (EVALUATION PRESERVES COST ANNOTATION CONTEXT VALIDITY). *If $\tilde{\Gamma}; C \vdash M \Downarrow^A C'; \langle \tilde{v}, c \rangle$ and C is valid for $\tilde{\Gamma}$, then C' is valid for $\tilde{\Gamma}$ and for \tilde{v} .*

Again, we only consider valid cost annotation contexts in what follows.

4.3 Clairvoyance evaluation cost

The total evaluation cost under the Annotated Clairvoyance Semantics can be recovered by adding up each individual cost as appeared in the annotations given to the thunks created during this evaluation, as well as the cost annotation given to the output value. We define the following operation for summing over cost annotations from a cost annotation context C :

$$\text{sumcost}_A(C) := \sum_{c \in \text{im}(C)} c$$

For an evaluation $\tilde{\Gamma}; C \vdash M \Downarrow^A C'; \langle \tilde{v}, c \rangle$, the cost incurred by the evaluation can be computed by

$$c + \text{sumcost}_A(C' \setminus C)$$

Like before, the difference $C' \setminus C$ removes all thunks and their annotations existing in the initial context, as we do not want to take into account any cost due to computations prior to the current one. Depending on the nondeterministic choices made during the evaluation, this total cost may be the call-by-need cost or, instead, a more eager cost.

Correspondence between the two Clairvoyance Semantics. The Annotated Clairvoyance Semantics is sound and complete with respect to the original one, and the corresponding evaluations produce exactly the same cost. To prove this claim, we first define a correspondence between values and environments of the two semantics.

Definition 4.4 (Corresponding Clairvoyance values and environments). Let \tilde{v} be a value of the Annotated Clairvoyance Semantics and \hat{v} a value of the monadic Clairvoyance Semantics. They are *corresponding values* modulo names, denoted $\tilde{v} \sim \hat{v}$, if the rules in Figure 8 apply. This relation extends to environments naturally.

$$\begin{array}{c}
 \frac{t \in \{\text{true}, \text{false}, \text{nil}, \perp\}}{t \sim t} \quad \frac{\tilde{v}_1 \sim \hat{v}_1 \quad \tilde{v}_2 \sim \hat{v}_2}{\text{cons } \tilde{v}_1 \tilde{v}_2 \sim \text{cons } \hat{v}_1 \hat{v}_2} \quad \frac{\tilde{v}_1 \sim \hat{v}_1 \quad \tilde{v}_2 \sim \hat{v}_2}{(\tilde{v}_1, \tilde{v}_2) \sim (\hat{v}_1, \hat{v}_2)} \quad \frac{}{\tilde{v} \sim \hat{v}}
 \end{array}$$

Fig. 8. The value relation \sim between annotated clairvoyance and original clairvoyance values.

If \tilde{v} and \hat{v} are related by \sim , the latter can be regarded as effectively abstracting away the name from the former. Furthermore, a value \hat{v} in the monadic Clairvoyance Semantics corresponds to an infinite set of values \tilde{v} in the Annotated Clairvoyance Semantics that are all structurally the same but with different names to thunks, if the values have thunks nested inside.

The theorems below establish the soundness and completeness together with cost equality. In these two theorems we consider the evaluation of a well-typed term $\gamma : M \vdash A$.

THEOREM 4.5 (SOUNDNESS OF THE ANNOTATED CLAIRVOYANCE SEMANTICS WRT THE MONADIC CLAIRVOYANCE SEMANTICS). *Let $\tilde{\Gamma}; C \vdash M \Downarrow^A C'; \langle \tilde{v}, c \rangle$, with the cost annotation context C valid for $\tilde{\Gamma}$. Then for all $\hat{\Gamma} \in [\![\gamma]\!]$ with $\tilde{\Gamma} \sim \hat{\Gamma}$, there exists $(\hat{v}, c') \in [\![M]\!](\hat{\Gamma})$, such that $\tilde{v} \sim \hat{v}$ and $c + \text{sumcost}_A(C' \setminus C) = c'$.*

THEOREM 4.6 (COMPLETENESS OF THE ANNOTATED CLAIRVOYANCE SEMANTICS WRT THE MONADIC CLAIRVOYANCE SEMANTICS). *Let $\hat{\Gamma} \in [\![\gamma]\!]$ with $\tilde{\Gamma} \sim \hat{\Gamma}$ and let C be a cost annotation context valid for $\tilde{\Gamma}$. Then for all $(\hat{v}, c) \in [\![M]\!](\hat{\Gamma})$, there is an evaluation $\tilde{\Gamma}; C \vdash M \Downarrow^A C'; \langle \tilde{v}, c' \rangle$ with $\tilde{v} \sim \hat{v}$ and $c' + \text{sumcost}_A(C' \setminus C) = c$.*

4.4 Correspondence between the Memorist and the Annotated Clairvoyance Semantics

Clairvoyance evaluation can make nondeterministic choices that are more eager than necessary. Thus in principle, to show that the Memorist Semantics indeed corresponds to the lazy semantics, we need to choose the right nondeterministic branch of Clairvoyance evaluation for a particular demand, which is difficult to do directly.

In our work, we instead take an approach inspired by the cost minimality and existence proofs of Xia et al. [2024]. That is, we prove that the cost inferred from the Memorist Semantics is never larger than any nondeterministic branch of a corresponding Clairvoyance evaluation, and that there is always a Clairvoyance branch that produces the same cost. Together they imply that the inferred cost is exactly the minimal cost that can be successfully produced on any branch of a corresponding Clairvoyance evaluation, that is, the inferred cost is indeed the correct lazy cost. However, unlike their Demand Semantics which can be viewed as backwardly inferring the demand on input from that on the output to guide the nondeterministic evaluation choices, the Memorist Semantics instead provides a summary of all individual pieces of information at the end which we can cherry-pick according to the demand of interest and derive thunk usage and cost from them. As a result, we cannot simply sum up and compare the evaluation cost step by step, since essential information including thunk usage is generally not available at the point a thunk is created (which is when it would be already evaluated or skipped in a Clairvoyance evaluation).

To address this problem we prove first that the thunk usage tracked by the Memorist semantics is both *minimal* and *sufficient* with respect to the Clairvoyance Semantics. By minimality, we mean the inferred usage set is never more than what is actually evaluated in any corresponding Clairvoyance evaluation. By sufficiency, we assert there is some corresponding Clairvoyance evaluation that evaluates exactly those thunks as recorded by the Memorist Semantics. By showing the usage captures exactly the lazy behaviour, we can then straightforwardly derive the cost correctness.

$$\begin{array}{c}
 \frac{t \in \{\text{true}, \text{false}, \text{nil}\}}{t \sim_f t} \quad \frac{\tilde{v}_1 \sim_f v_1 \quad \tilde{v}_2 \sim_f v_2}{\text{cons } \tilde{v}_1 \tilde{v}_2 \sim_f \text{cons } v_1 v_2} \quad \frac{\tilde{v}_1 \sim_f v_1 \quad \tilde{v}_2 \sim_f v_2}{(\tilde{v}_1, \tilde{v}_2) \sim_f (v_1, v_2)} \\
 \\
 \frac{\tilde{v} \sim_f v \quad f(i) = j}{\text{thunk}_i \tilde{v} \sim_f \text{thunk}_j v} \quad \frac{}{\perp \sim_f \text{thunk}_j v}
 \end{array}$$

Fig. 9. The relation \sim_f between Memorist and Annotated Clairvoyant values using renaming function f .

Renaming. Corresponding thunks in Memorist Semantics and Annotated Clairvoyance may have different names. To address this discrepancy, we define renaming functions of type $\mathcal{N}_A \rightarrow \mathcal{N}_M$ that rename Annotated Clairvoyance thunk names $\mathcal{N}_A \subset \mathbb{N}$ to Memorist thunk names $\mathcal{N}_M \subset \mathbb{N}$. Since the Memorist evaluation is always eager, it will never evaluate fewer thunks than the corresponding Annotated Clairvoyance evaluation. Hence, thunk renaming functions are always total.

We define a notion of validity for renaming functions to ensure that the names in the domains and images are indeed names assigned in the respective evaluation. A renaming function is also required to be injective so that two names, if related under this function, are uniquely related.

Definition 4.7 (Validity of renaming functions). A renaming function $f : \mathcal{N}_A \rightarrow \mathcal{N}_M$ is *valid* with respect to some valid Memorist annotation context C and valid Clairvoyance cost annotation context \mathcal{A} if f is injective with $\text{dom}(f) = \text{dom}(C)$ and $\text{im}(f) \subseteq \text{dom}(\mathcal{A})$.

4.5 Functional Correctness of Memorist Semantics

We define a relation $\tilde{v} \sim_f v$ on an annotated Clairvoyance value \tilde{v} and a Memorist value v , parametrised by a renaming function $f : \mathcal{N}_A \rightarrow \mathcal{N}_M$.

Definition 4.8 (Value-name and environment-name correspondence). Given a renaming function f as above, value-name correspondence $\tilde{v} \sim_f v$ is defined inductively by the rules in Fig. 9. Environment-name correspondence $\tilde{\Gamma} \sim_f \Gamma$ for an annotated Clairvoyance environment $\tilde{\Gamma}$ and a Memorist environment Γ holds if $\text{dom}(\tilde{\Gamma}) = \text{dom}(\Gamma)$ and $\forall x \in \text{dom}(\Gamma), \tilde{\Gamma}(x) \sim_f \Gamma(x)$.

The relation \sim_f describes the partial correspondence on evaluated values/environments and names between the two semantics. Given two environments related by \sim_f through a valid renaming function f , a Memorist evaluation and an Annotated Clairvoyance evaluation of the same term produce \sim_f -related values, and the extended renaming function remains valid. Validity of the renaming functions ensure that each evaluated thunk in the Clairvoyance evaluation corresponds to a unique thunk name from the Memorist evaluation.

Similarly, we define a thunkwise cost correspondence:

Definition 4.9 (Thunkwise cost correspondence). Let \mathcal{A} and C be a Memorist and a Clairvoyance annotation context respectively, and f a valid thunk renaming function with respect to \mathcal{A}, C . \mathcal{A} is *thunkwise cost corresponded* with respect to C under f , denoted $\mathcal{A} \dot{\sim}_f C$, if $\forall i \in \text{dom}(C), C(i) = \pi_1(\mathcal{A}(f(i)))$.

The theorem below guarantees that, for each thunk evaluated in a Clairvoyance evaluation, there is a thunk evaluated in the corresponding Memorist evaluation with their names related by a thunk renaming function and their cost annotations being equal. Intuitively, the equality holds because any potential difference in evaluation cost is due to the Clairvoyance evaluation nondeterministically evaluating or skipping a thunk, while a cost annotation to a value records only the portion of cost incurred by an evaluation to that value that is never thunked before.

$$\begin{array}{c}
 \frac{t \in \{\text{true}, \text{false}, \text{nil}\}}{t \leq t} \quad \frac{\hat{v}_1 \leq v_1 \quad \hat{v}_2 \leq v_2}{\text{cons } \hat{v}_1 \hat{v}_2 \leq \text{cons } v_1 v_2} \quad \frac{\hat{v}_1 \leq v_1 \quad \hat{v}_2 \leq v_2}{(\hat{v}_1, \hat{v}_2) \leq (v_1, v_2)} \\
 \\
 \frac{\hat{v} \leq v}{\text{thunk } \hat{v} \leq \text{thunk}_i v} \quad \frac{}{\perp \leq \text{thunk}_i v}
 \end{array}$$

Fig. 10. Definition of \leq between an unnamed partially evaluated Clairvoyant value and a Memorist value

THEOREM 4.10 (FUNCTIONAL AND COST ANNOTATION CORRECTNESS). Let $\Gamma; \mathcal{A} \vdash M \Downarrow \mathcal{A}'; \langle v, (c, s) \rangle$ and $\tilde{\Gamma}; C \vdash M \Downarrow^A C'; \langle \tilde{v}, c' \rangle$, and f be a valid renaming function wrt \mathcal{A} and C . If $\tilde{\Gamma} \sim_f \Gamma$ and $\mathcal{A} \sim_f \mathcal{A}'$, then $\tilde{v} \sim_f v$, $c = c'$ and $C' \sim_f \mathcal{A}'$, for some f' extending f and valid wrt \mathcal{A}' and C' .

4.5.1 Deriving usage based on demands. What needs to be evaluated in a lazy evaluation is driven by demand which cannot always be determined locally. The behaviour of Clairvoyance evaluation, in particular, varies according to such external demand. In Memorist Semantics, we consider evaluation with respect to demands on output via usage sets.

In the Memorist Semantics, each of the thunked and final values is annotated with a cost and a usage set. This usage set contains names of thunks needed for the evaluation to reach the outermost portion of the annotated value. As discussed previously, if more of the value is demanded, we can compute a set of names of all thunks that are needed with respect to this external demand, by collecting altogether (1) all the thunks in the value according to the demand, (2) the thunks in their usage set annotations, and (3) thunks in the usage set annotated to the output value.

Since the demand is provided externally and we do not generally know what names are given to a thunk from the outset, we need a way to abstract away from names when representing demand. For the current proof, the partially evaluated values (with unnamed thunks) defined for the monadic Clairvoyance Semantics in Section 4.1 provides a convenient way in this regards. To distinguish between the values from the two Clairvoyance Semantics in what follows, we sometimes refer to the values for the monadic Clairvoyance Semantics as *thunk-unnamed* partially evaluated values and those for the annotated variant as *thunk-named* partially evaluated values. We define the following relation to relate thunk-unnamed values with the Memorist values.

Definition 4.11 (Partially evaluated values and environments). Given a Memorist value v and a thunk-unnamed partially evaluated value \hat{v} of the same type, \hat{v} is a partially evaluated version of v , denoted by $\hat{v} \leq v$, if the rules in Fig. 10 apply. The relation extends naturally to environments defined for the same type context.

Thus, for a Memorist value v , we use \hat{v} with $\hat{v} \leq v$ to represent a demand on v . The following relation characterises the set of names of all thunks needed to be evaluated given such a demand. It is easy to show that, such a set always exists for any external demand $\hat{v} \leq v$, and that a \hat{v} satisfying \leq always exists given v and a set s of thunk names.

Definition 4.12 (Usage representation sets). Given a Memorist value v , a valid annotation context \mathcal{A} , and a partially evaluated value \hat{v} with $\hat{v} \leq v$, the relation $\text{Repr}_{\mathcal{A}}(v, \hat{v}, s)$ is defined by the rules in Figure 11. The relation extends to environments in the following way. Given a Memorist environment Γ and a partially evaluated environment $\hat{\Gamma}$ defined for the same type context with $\hat{\Gamma} \leq \Gamma$, $\text{Repr}_{\mathcal{A}}(\Gamma, \hat{\Gamma}, s)$ if

$$s = \bigcup_{x \in \text{dom}(\Gamma)} \{i \in s' \mid \text{Repr}_{\mathcal{A}}(\Gamma(x), \hat{\Gamma}(x), s')\}$$

$$\begin{array}{c}
 \frac{\begin{array}{c} t \in \{\text{true}, \text{false}, \text{nil}\} \\ \text{Repr}_{\mathcal{A}}(t, t, \emptyset) \end{array}}{\text{Repr}_{\mathcal{A}}(v_1, \hat{v}_1, s_1)} \quad \frac{\text{Repr}_{\mathcal{A}}(v_1, \hat{v}_1, s_1) \quad \text{Repr}_{\mathcal{A}}(v_2, \hat{v}_2, s_2)}{\text{Repr}_{\mathcal{A}}(\text{cons } v_1 v_2, \text{cons } \hat{v}_1 \hat{v}_2, s_1 \cup s_2)} \quad \frac{}{\text{Repr}_{\mathcal{A}}(\text{thunk}_i v, \perp, \emptyset)} \\
 \frac{\text{Repr}_{\mathcal{A}}(v_1, \hat{v}_1, s_1) \quad \text{Repr}_{\mathcal{A}}(v_2, \hat{v}_2, s_2)}{\text{Repr}_{\mathcal{A}}((v_1, v_2), (\hat{v}_1, \hat{v}_2), s_1 \cup s_2)} \quad \frac{\text{Repr}_{\mathcal{A}}(v, \hat{v}, s) \quad \pi_2(\mathcal{A}(i)) = s'}{\text{Repr}_{\mathcal{A}}(\text{thunk}_i v, \text{thunk } \hat{v}, \{i\} \cup s \cup s')}
 \end{array}$$

Fig. 11. Definition of the relation $\text{Repr}_{\mathcal{A}}$, parametrised on an annotation context \mathcal{A} , between a Memorist value v , a demand \hat{v} on v , and a set s of thunk names.

Furthermore, let (c, s') be the annotation given to v . We call the set $s \cup s'$ the *usage representation set* for the demand \hat{v} on v .

LEMMA 4.13. *Let v and \hat{v} be a Memorist value and a thunk-unnamed partially evaluated value \hat{v} respectively, of the same type. Let \mathcal{A} be a valid annotation context for v . We have $\hat{v} \leq v$ if and only if there is a set s of thunk names with $\text{Repr}_{\mathcal{A}}(v, \hat{v}, s)$. The same holds for environments in both directions.*

4.5.2 Usage minimality. As previously mentioned, cost correctness will be proved via the correctness of thunk usage. The latter is to be established in two parts, usage minimality and usage sufficiency, that together characterise the correspondence with the laziest possible Clairvoyance evaluation. We first prove that the usage representation set inferred from the Memorist Semantics is minimal, in the sense that it is never more than what is actually evaluated in any corresponding successful Clairvoyance evaluation. More specifically, we show that every name in such a usage representation set always corresponds to the name of an evaluated thunk in a corresponding successful Annotated Clairvoyance evaluation, as related by some valid thunk renaming function.

One observation can immediately be made. If the minimality in the above sense indeed holds between a Memorist evaluation and a corresponding annotated Clairvoyance evaluation, it will be the case that, for every Memorist thunk j that corresponds to an evaluated thunk in the Clairvoyance evaluation (with their names related by a thunk renaming function f), its annotation usage set must always be a subset of the image of f , $\text{im}(f)$. Recall that the $\text{im}(f)$ is the set of names of all Memorist thunks that indeed correspond to some Clairvoyance thunks. This is desired; otherwise we would have discovered a thunk inferred to be needed by the Memorist evaluation whereas the successful Clairvoyance evaluation did not actually evaluate the thunk, failing minimality in the above sense. We define a relation to capture this observation.

Definition 4.14 (Thunkwise usage minimal). A Memorist annotation context \mathcal{A} is *thunkwise usage minimal* with respect to a Clairvoyance annotation context C under a valid thunk renaming function f , denoted $\mathcal{A} \sqsubseteq_f C$, if $\forall i \in \text{dom}(C)$, we have $\pi_2(\mathcal{A}(f(i))) \subseteq \text{im}(f)$.

The validity of f ensures that $i \in \text{dom}(f)$ and $f(i) \in \text{dom}(\mathcal{A})$. The following lemma states that any two corresponding evaluations of some well-typed term $\gamma \vdash M : A$ preserve the thunkwise usage minimality in the above sense. In addition, the usage representation set in the annotation to the output value also satisfies a similar requirement: it is a subset of the image of the extended renaming function.

LEMMA 4.15 (THUNKWISE USAGE MINIMALITY). *Let $\Gamma ; \mathcal{A} \vdash M \Downarrow \mathcal{A}' ; \langle v, (c, s) \rangle$ and $\tilde{\Gamma} ; C \vdash M \Downarrow^A C' ; \langle \tilde{v}, c' \rangle$, and let f be a valid thunk renaming function. If $\tilde{\Gamma} \sim_f \Gamma$, $\mathcal{A} \dot{\sim}_f C$ and $\mathcal{A} \sqsubseteq_f C$, then $\mathcal{A}' \sqsubseteq_{f'} C'$ and $s \subseteq \text{im}(f')$, as well as $\tilde{v} \sim_{f'} v$ and $C' \dot{\sim}_f \mathcal{A}'$, for some valid renaming function f' extending f .*

From here we can establish the desired usage minimality, generalising to the usage representation set with respect to an arbitrary external demand on the output value. The notation $f[\cdot]$ denotes the image of a set under some function f .

THEOREM 4.16 (USAGE MINIMALITY). *Let $\Gamma; \mathcal{A} \vdash M \Downarrow \mathcal{A}'; \langle v, (c, s) \rangle$ and $\tilde{\Gamma}; C \vdash M \Downarrow^A C'; \langle \tilde{v}, c' \rangle$, and f be a valid thunk renaming function, with $\tilde{\Gamma} \sim_f \Gamma$, $\mathcal{A} \sim_f C$ and $\mathcal{A} \Subset_f C$. Apart from satisfying the thunkwise usage minimality as stated above for some valid renaming function f' extending f , the following also holds:*

- Given any thunk-unnamed partially evaluated value \hat{v} with $\tilde{v} \sim \hat{v}$ and $\text{Repr}_{\mathcal{A}'}(v, \hat{v}, s')$, we have $s \cup s' \subseteq f'[\text{dom}(C')]$, and moreover, $(s \cup s') \setminus \text{dom}(\mathcal{A}) \subseteq f'[\text{dom}(C' \setminus C)]$.

The additional conclusion stated with set difference allows us to excluding thunks already exist prior to this evaluation. In short, the theorem states that thunks captured by $(s \cup s') \setminus \text{dom}(\mathcal{A})$, i.e., all thunks created during the evaluation in question and inferred to be needed by the Memorist Semantics through usage tracking, always correspond (under thunk renaming f') to a subset of all thunks evaluated on any successful branch in the corresponding Clairvoyance evaluation, where the latter is captured by $\text{dom}(C' \setminus C)$. That is, a corresponding Clairvoyance evaluation may evaluate more thunks, but it can never evaluated less. Thus, the inferred usage by the Memorist Semantics is minimal in this sense.

4.5.3 Usage sufficiency. Usage minimality essentially states that the usage set inferred from the Memorist semantics is ‘not too big’. Next, we prove usage sufficiency, which states the inferred usage set is ‘not too small’: for any Memorist evaluation, there is a corresponding Annotated Clairvoyance evaluation such that what is evaluated in the latter corresponds exactly to the usage set inferred by the former.

In general, we cannot assert there is always a successful Clairvoyance evaluation given an arbitrary environment $\tilde{\Gamma}$ even if it satisfies $\tilde{\Gamma} \sim_f \Gamma$ for some Memorist evaluation environment Γ and a valid renaming function f . For instance, evaluating force x in an environment $\tilde{\Gamma}$ that maps x to a skipped thunk \perp can never succeed. To address this, we utilise the usage representation set, which supposedly contains all thunks needed to be evaluated, to infer a minimally workable environment. We have shown that this set is minimal in that every thunk in it must be evaluated in all corresponding successful Clairvoyance evaluation. It remains a question whether it might be ‘too minimal’ such that it misses some thunks that should have been evaluated, though it does not matter for now since we are not asserting anything new about the derived environment. We only rely on it to set a ‘lower bound’ on how less evaluated the Clairvoyance environment can be.

To do so, we first observe that, given the union of a usage representation set and the output usage annotation set, if we take the intersection of this union and the domain of the initial annotation context \mathcal{A} (where the latter contains all thunks that already exist before the evaluation), the resulting set s should contain all existing thunks that are inferred to be needed by the evaluation of interest. From s we can construct a partially evaluated environment $\hat{\Gamma}$ satisfying $\text{Repr}_{\mathcal{A}}(\Gamma, \hat{\Gamma}, s)$, which can be viewed as the minimal demand on the environment. We can then consider any Clairvoyance environment that is related to the given Memorist environment under \sim_f and no less eagerly evaluated than the minimum environment $\hat{\Gamma}$.

We define a relation to express what it means for an Annotated Clairvoyance value or environment to be no less eagerly evaluated as a given demand.

Definition 4.17 (No less eagerly evaluated values and environments). Let \tilde{v} and \hat{v} be an Annotated Clairvoyance value and a thunk-unnamed partially evaluated value respectively of the same type. \tilde{v} is no less eagerly evaluated than \hat{v} , denoted $\hat{v} \lesssim \tilde{v}$, if the rules in Figure 12 apply.

$$\begin{array}{c}
 981 \quad t \in \{\text{true}, \text{false}, \text{nil}\} \\
 982 \quad \frac{}{t \lesssim t} \qquad \hat{v}_1 \lesssim \tilde{v}_1 \quad \hat{v}_2 \lesssim \tilde{v}_2 \\
 983 \quad \text{cons } \hat{v}_1 \hat{v}_2 \lesssim \text{cons } \tilde{v}_1 \tilde{v}_2 \qquad \frac{\hat{v}_1 \lesssim \tilde{v}_1 \quad \hat{v}_2 \lesssim \tilde{v}_2}{(\hat{v}_1, \hat{v}_2) \lesssim (\tilde{v}_1, \tilde{v}_2)} \\
 984 \\
 985 \quad \frac{\hat{v} \lesssim \tilde{v}}{\text{thunk } \hat{v} \lesssim \text{thunk}_i \tilde{v}} \qquad \frac{}{\perp \lesssim \text{thunk}_i \tilde{v}}
 \end{array}$$

Fig. 12. Definition of \lesssim between a thunk-unnamed partially evaluated value and a named one (i.e. a value in the Annotated Clairvoyance Semantics) of the same type.

The relation extends to environments naturally: let $\tilde{\Gamma}$ and $\hat{\Gamma}$ be an Annotated Clairvoyance environment and a thunk-unnamed partially evaluated environment respectively, defined on the same typing context. We say $\hat{\Gamma} \lesssim \tilde{\Gamma}$ if $\forall x \in \text{dom}(\hat{\Gamma})$, $\hat{\Gamma}(x) \lesssim \tilde{\Gamma}(x)$.

We can now state and prove usage sufficiency. Below we consider a well-typed term $\gamma \vdash M : A$.

THEOREM 4.18 (USAGE SUFFICIENCY). *Let $\Gamma; \mathcal{A} \vdash M \Downarrow \mathcal{A}' ; \langle v, (c, s) \rangle$ with \mathcal{A} a valid annotation context for Γ . Let $\hat{v} \in \llbracket A \rrbracket$ with $\text{Repr}_{\mathcal{A}'}(v, \hat{v}, s')$ for some set s' of thunk names. There is some $\tilde{\Gamma} \in \llbracket \mathcal{Y} \rrbracket$ satisfying $\text{Repr}_{\mathcal{A}}(\Gamma, \tilde{\Gamma}, (s \cup s') \cap \text{dom}(\mathcal{A}))$, such that for all Annotated Clairvoyance environment $\tilde{\Gamma}$ with $\hat{\Gamma} \lesssim \tilde{\Gamma}$, a valid annotation context C for $\tilde{\Gamma}$ and a valid thunk renaming function f with $\tilde{\Gamma} \sim_f \Gamma$, the following holds.*

- If f, \mathcal{A} and C satisfy the assumptions in the demand minimality theorem, then there is some $\tilde{\Gamma}; C \vdash M \Downarrow^A C' ; \langle \tilde{v}, c' \rangle$ with $\hat{v} \lesssim \tilde{v}$ and $(s \cup s') \setminus \text{dom}(\mathcal{A}) = f[\text{dom}(C' \setminus C)]$, where f' is some valid renaming function extending f .

The theorem intuitively states that, given a Memorist evaluation with some demand, we can always construct a corresponding Clairvoyance evaluation, such that all the thunks created during the Memorist evaluation and inferred to be actually needed (as captured by the set $(s \cup s') \setminus \text{dom}(\mathcal{A})$) coincide with all the thunks evaluated within the corresponding Clairvoyance evaluation (as captured by the set $\text{dom}(C' \setminus C)$), when thunk renaming is accounted for. In other words, there is a always a corresponding Clairvoyance evaluation that evaluates exactly as much as the inferred usage from the Memorist evaluation. Hence, we say the inferred usage is sufficient.

Together, the usage minimality and sufficiency theorems imply that the inferred lazy thunk usage from the Memorist Semantics is minimal but nontrivial: it is never more than what will be evaluated in any corresponding successful Clairvoyance evaluation, while indeed corresponding to some successful Clairvoyance evaluation. To conclude, the inferred usage represents exactly what is evaluated on the laziest branch in a corresponding successful Clairvoyance evaluation.

4.6 Lazy cost correspondence

In the Memorist Semantics, lazy evaluation cost is derived based on the inferred evaluation usage and thunks' annotated cost. We have shown that, the individual cost annotations are correct (Theorem 4.10), and the inferred thunk usage corresponds to what is actually evaluated in the corresponding laziest possible Clairvoyance evaluation (Theorems 4.16 and 4.18). Consequently, the derived lazy cost from the Memorist Semantics coincides with the evaluation cost of the laziest possible Clairvoyance evaluation.

Again, since it is not straightforward to characterise the lazy Clairvoyance evaluation branch directly, we prove instead the following *cost minimality* and *cost existence* results. We consider evaluation of a well-typed term $\gamma \vdash M : A$.

1030 THEOREM 4.19 (COST MINIMALITY). Let $\Gamma; \mathcal{A} \vdash M \Downarrow \mathcal{A}'; \langle v, (c, s) \rangle$ and $\tilde{\Gamma}; C \vdash M \Downarrow^A C'; \langle \tilde{v}, c' \rangle$,
 1031 and f be a valid thunk renaming function. Let $\tilde{\Gamma} \sim_f \Gamma$, $\mathcal{A} \sim_f C$ and $\mathcal{A} \Subset_f C$.

1032 Then given any thunk-unnamed partially evaluated value \hat{v} with $\tilde{v} \sim \hat{v}$ and $\text{Repr}_{\mathcal{A}'}(v, \hat{v}, s')$, we have

$$1033 \quad c + \text{sumcost}_M(\mathcal{A}' \setminus \mathcal{A}, s \cup s') \leq c' + \text{sumcost}_A(C' \setminus C).$$

1035 THEOREM 4.20 (COST EXISTENCE). Let $\Gamma; \mathcal{A} \vdash M \Downarrow \mathcal{A}'; \langle v, (c, s) \rangle$ with \mathcal{A} a valid annotation context
 1036 for Γ . Let $\hat{v} \in \llbracket A \rrbracket$ with $\text{Repr}_{\mathcal{A}'}(v, \hat{v}, s')$ for some set s' of thunk names. There is some $\hat{\Gamma} \in \llbracket \mathcal{Y} \rrbracket$ satisfying
 1037 $\text{Repr}_{\mathcal{A}}(\Gamma, \hat{\Gamma}, (s \cup s') \cap \text{dom}(\mathcal{A}))$, such that for all Annotated Clairvoyance environment $\tilde{\Gamma}$ with $\hat{\Gamma} \lesssim \tilde{\Gamma}$,
 1038 a valid annotation context C for $\tilde{\Gamma}$ and a valid thunk renaming function f with $\tilde{\Gamma} \sim_f \Gamma$.

- If f , \mathcal{A} and C satisfy the assumptions in the usage minimality theorem, then there is some
 1040 $\hat{\Gamma}; C \vdash M \Downarrow^A C'; \langle \tilde{v}, c' \rangle$ with $\hat{v} \lesssim \tilde{v}$, and

$$1042 \quad c + \text{sumcost}_M(\mathcal{A}' \setminus \mathcal{A}, s \cup s') = c' + \text{sumcost}_A(C' \setminus C).$$

1043 In essence, cost minimality states that the derived lazy cost from the Memorist evaluation is
 1044 no larger than the cost incurred by any corresponding Clairvoyance evaluation. Cost existence
 1045 states that there is some corresponding Clairvoyance evaluation whose evaluation cost equals to
 1046 the derived lazy cost from the Memorist evaluation. We thus conclude that the Memorist Semantics
 1047 indeed derives the correct lazy evaluation cost.
 1048

1049 5 Implementation

1050 In this section, we provide a proof of concept implementation of the Memorist Semantics in Rocq
 1051 Prover. Our implementation uses shallow embedding and encapsulates all the operations over
 1052 thunks and annotations using a simple monadic interface.
 1053

1054 First, we define datatypes for representing thunks and annotations:

```
1055 Record Th (A : Type) : Type := MkTh { name : nat; val : A }.
1056 Definition Annot : Type := nat * NatSet.
```

1057 We model named thunks using the above `Th` type. An annotation is a pair of a natural number (`nat`)
 1058 and a finite set of natural number (`NatSet`), representing the cost and the usage set, respectively.
 1059

1060 An inductive type like `lists` is encoded as follows:

```
1061 Inductive ListT (A : Type) : Type :=
1062 | NilT : ListT A
1063 | ConsT : Th A -> Th (ListT A) -> ListT A.
```

1064 We implement the Memorist evaluation using a monad `M`, encapsulating the manipulations of
 1065 thunks and the accumulation of cost. The definition of `M` is as follows:
 1066

```
1067 Record Result (A : Type) : Type := MkRes
1068 { val : A; annot : Annot; cont : AC; annotC := fst annot; annotU := snd annot }.
1069 Notation "<[v, a, ac]>" := ( MkRes v a ac ).
```

```
1070 Definition M (A : Type) : Type := AC -> Result A.
```

1072 It represents the evaluation of an expression on an initial annotation context (`AC`) to a result (`Result`).
 1073 The annotation context is encoded using an association list. The result is a record containing the
 1074 result value (`val`), the annotation of the result (`annot`), and the final annotation context (`cont`). In
 1075 addition, we also use `annotC` and `annotU` to represent the all thunk names and all the usage sets of
 1076 `cont`, respectively. The monad `M` is essentially a generalized state monad, also known as an *update*
 1077 *monad* [Ahman and Uustalu 2013].
 1078

```

1079 Definition ret {A} (x : A) : M A := fun ac => <[x, (0,emptyset), ac]>.
1080
1081 Definition bind {A B} (m : M A) (f : A -> M B) : M B := fun ac =>
1082   let '<[x, a1, ac1]>' := m ac in let '<[y, a2, ac2]>' := (f x) ac1 in
1083     <[y, (fst a1 + fst a2, snd a1 ∪ snd a2), ac2]>.
1084 Notation "x >> f" := (bind x f).
1085 Notation "x >> y" := (bind x (fun _ => y)).
1086
1087 Definition lazy {A} (m : M A) : M (Th A) := fun ac =>
1088   let '<[x, a, ac1]>' := m ac in <[MkTh (nextIdx ac1) x, (0,emptyset), ext ac1 a]>.
1089 Notation "'lazy_letM' x := y 'in' z" := (bind (lazy y) (fun x => z)).
1090
1091 Definition collect (i : nat) : M unit := fun ac =>
1092   <[tt, (0, \{i\} ∪ mapS ac i), ac]>.
1093
1094 Definition forcing {A B} (th : Th A) (f : A -> M B) : M B :=
1095   match th with MkTh i v => collect i >> f v end.
1096 Notation "f $! x" := (forcing x f).
1097
1098 Definition force {A} (th : Th A) : M A := forcing th ret.
1099
1100 Definition tick : M unit := fun ac => <[tt, (1,emptyset), ac]>.
1101

```

1102 Fig. 13. Definitions of the monad operations and other basic operations of the metalanguage. All level and
 1103 associativity declarations for notations are omitted. Given a annotation context ac, the function nextIdx
 1104 ac returns a name not in the domain of ac, and ext ac a extends ac by mapping the next new name (as
 1105 will be returned by nextIdx) to the annotation a. The expression mapS ac i maps a name i to its usage set
 1106 annotation under ac. The notation \{i\} denotes a singleton set containing i.

1107
 1108
 1109
 1110 All the operations over thunks and annotations in the Memorist Semantics can be encapsulated
 1111 using a few combinators of the M monad, shown in Fig. 13. The ret (for return) and bind functions
 1112 are standard monadic combinators. The lazy function wraps the value in M with a thunk; the
 1113 forcing and force functions unwrap thunks. Finally, the tick function increases the cost by 1. The
 1114 set of combinators is the same as that of the Clairvoyance Monad [Li et al. 2021].

1115 To analyse a program, we can use the operational semantics of the Memorist Semantics (Fig. 6)
 1116 as a recipe to translate a pure program to a monadic program that *reifies* the cost. Let's consider
 1117 again the functions described in Section 2. We show the translation of those functions in Fig. 14.

1118 We affix the names of the translated functions with an M to distinguish them from the pure
 1119 original version. The translated program operates on thunked lists. Its structure follows closely that
 1120 of the original definition. Function and constructor arguments are generally thunked. Whenever a
 1121 function needs to apply to an argument, the thunks need to be unwrapped first, which is done
 1122 via forcing (\$!). We separate the transformed definition of append into a top-level appendM and an
 1123 auxiliary appendM_ (similar for truePrefix) so they can pass Rocq Prover's termination checker
 1124 directly. Since the cost model employed here is the number of function calls, we put a tick at the
 1125 start of each call to increment cost by one. Note that we do not put ticks in the top-level definitions
 1126 appendM and truePrefixM, as they are simply unwrapping the thunks around the arguments.

```

1128 Fixpoint appendM_ {A} (xs : ListT A) (tys : Th (ListT A)) : M (ListT A) :=
1129   tick >>
1130   match xs with
1131   | NilT => force tys
1132   | ConsT tx txs =>
1133     lazy_letM tzs := (fun xs' => appendM_ xs' tys) $! txs in
1134     ret (ConsT tx tzs)
1135   end.
1136
1137 Definition appendM {A} (txs tys : Th (ListT A)) : M (ListT A) :=
1138   (fun xs => appendM_ xs tzs) $! txs.
1139
1140 Fixpoint truePrefixM_ (xs : ListT bool) : M (ListT bool) :=
1141   tick >>
1142   match xs with
1143   | NilT => ret NilT
1144   | ConsT tx txs =>
1145     lazy_letM tzs := truePrefixM_ $! txs in
1146     (fun x => if x then ret (ConstT tx tzs) else ret NilT) $! tx
1147   end.
1148
1149 Definition truePrefixM (txs : Th (ListT bool)) : M (ListT bool) :=
1150   truePrefixM_ $! txs.
1151
1152 Definition truePrefixOfAppendM (txs : Th (ListT bool)) (tys : Th (ListT bool))
1153   : M (ListT bool) := tick >>
1154   lazy_letM tzs := appendM txs tys in truePrefixM tzs.
1155

```

Fig. 14. Translation of the functions `append`, `truePrefix` and `truePrefixOfAppend`.

To analyse cost, we still need some additional operations to collect thunks and summing up their cost. For the current example, we are interested in the lazy cost with respect to the length of a list demanded; accordingly, we can define the following functions to collect the names of thunks in the value that are demanded, along with their thunk usage.

```

1164 Fixpoint usageL_ (ac : AC) {A} (n : nat) (xs : ListT A) : NatSet :=
1165   match xs, n with
1166   | ConsT (MkTh i xs') (MkTh j xs'), S n' =>
1167     \{i\} \cup mapS ac i \cup \{j\} \cup mapS ac j \cup usageL_ n' xs' ac
1168   | _, _ => emptyset
1169   end.
1170 Definition usageL {A} (r : Result (ListT A)) (n : nat) : NatSet :=
1171   usageL_ (cont r) (val r) n.
1172

```

Here n is the number of thunked cons cells demanded in the output list. The notation $\{i\}$ denotes a singleton set containing i , and $\text{mapS } ac \ i$ maps the thunk name i to its usage set annotation. In general, different usage collection functions need to be defined for other datatypes and demands.

1173

As explained when introducing the formal semantics, to derive the lazy cost from the annotated result, we need to consider both thunks collected from the value according to the demand *and* the thunks in the output annotation. Moreover, we need to remove all thunks already present in the input environment. These can be realised using set union and difference. The output cost annotation must also be added in. We thus encapsulate this calculation into the following function. In the definition, the expression `dom ac` gives the domain of the annotation context `ac` as a `NatSet`, and `sumcost` sums over the cost annotated to the thunk names in the given set.

```
1177  Definition infcost {A} (r : Result A) (s : NatSet) (ac : AC) : nat :=
1178    let s1 := diff (annotU r ∪ s) (dom ac) in sumcost (cont r) s1 + annotC r
1186
```

Now let us consider a concrete example. Previously in [Section 2](#), we analysed informally the cost of the example program `truePrefix0fAppend [true;false] [true]`. Because the semantics is computable, we can obtain the concrete cost of this program by translating the program into the monadic language, running the translated program directly and applying `infcost` to derive the cost. In general, the input is a value obtained from some previous computations and thunked. For this concrete example, we simply lift them to the type `Th (ListT bool)`. This can be done automatically; more details can be found in the artefact. The following two are example transformations of the two source input lists respectively:

```
1195  MkTh 4 (ConsT (MkTh 3 true) (MkTh 2 (ConsT (MkTh 1 false) (MkTh 0 NilT))))
1196  MkTh 7 (ConsT (MkTh 6 true) (MkTh 5 NilT))
```

Below let `t11` and `t12` refer to these two lifted lists respectively and let `ac` refer to the annotation context produced by the transformation. This `ac` records the annotations of all the thunks in the lifted lists. Simply lifting the lists means each thunk so far is given the empty annotation `(0,emptyset)`, although this does not affect reasoning about the cost of `truePrefix0fAppendM` since these thunks, as part of the input, are removed via set difference when deriving the program cost.

The cost of the program can be inferred by applying `infcost` on the translated program and the set of thunk names collected according to some demand on the output. As before, we consider two different demands on the same output: one demanding only to the outermost constructor of the output list (*i.e.*, demanding 0 cons cells), and the other demanding one more cons cell—which happens to also be the last and only cons cells in the output.

```
1207  Compute let r := truePrefix0fAppendM t11 t12 ac in
1208    let cost0 := infcost r (usageL r 0) ac in
1209    let cost1 := infcost r (usageL r 1) ac in (cost0, cost1).
```

The above computes to $(3, 5)$, that is, the lazy cost with respect to the first demand is 3 while it is 5 for the second. We intentionally present it this way to highlight that, with the Memorist semantics, the program only needs to be evaluated once for analysing different demand.

The exact names of thunks and the annotations to thunks in the input lists may vary with different initial annotation contexts and different previous computations to constructing the input lists. Nevertheless, this will not change the inferred cost, as previous thunks are removed via set difference when inferring cost. Readers can find such examples of running the same program on different initial contexts in the artefact. Together with the semantics being computable, we contemplate that it may be possible to develop the semantics into a framework for testing.

More generally, we can also state and prove specifications of lazy cost for these functions. The first theorem below states that the lazy cost of `appendM` is linear in the length demanded from the output list but no more than the size of the first input list (since `append` never steps into the second input list), where the size here is measured using the function `sizeT` that counts the number of constructors in a list. The second states that the lazy cost of `truePrefixM` is linear in the

length demanded from the output. The AC0k and Name0k premises ensures the existing thunk names and annotations are valid in the sense described when previously when introducing the formal semantics. The Name0k predicate is defined for all types belonging to the TNamed class.

```
1229 Theorem appendM_cost :
1230   forall ac {A} (txs tys : Th (ListT A)) n (r := appendM txs tys ac),
1231     AC0k ac -> Name0k txs ac -> Name0k tys ac -> n < sizeT (val r) ->
1232       infcost r (usageL r n) ac = min (n + 1) (sizeT (Th.val txs)).
1233 
```

```
1234 Theorem truePrefixM_cost : forall ac txs n (r := truePrefixM txs ac),
1235   AC0k ac -> Name0k txs ac -> n < sizeT (val r) ->
1236     infcost r (usageL r n) ac = n + 1.
1237 
```

We can then prove the following cost specification for truePrefixOfAppendM, directly expressing that the lazy cost of truePrefixOfAppendM in terms of the cost of appendM and of truePrefixOfAppendM. The ‘plus 1’ at the end is due to the one additional cost incurred by truePrefixOfAppendM itself.

```
1241 Theorem truePrefixOfAppendM_cost :
1242   forall ac txs tys n d (r := truePrefixOfAppendM txs tys ac),
1243     AC0k ac -> Name0k txs ac -> Name0k tys ac -> n < sizeT (val r) ->
1244       infcost r (usageL r n) ac = min (n + 1) (sizeT (Th.val txs)) + (n + 1) + 1.
1245 
```

The proof, omitted here, makes direct use of the above theorems for appendM and truePrefixOfAppendM. We can also prove the following bound for truePrefixOfAppendM, especially when the size of the output list is smaller than that of the input list (with other conditions being the same):

```
1249   infcost r (usageL r n) ac <= 2 * sizeT (val r) + 1
```

This states that the cost is at most linear in the size of the output list. It is a somewhat loose bound and does not explicitly consider any demand on the final output. However, it reflects the inherent laziness in this program: when evaluating lazily, append need not step further after the first false is encountered in the (first) input lists, as it is where truePrefix stops processing the rest of the list.

6 Related Work

The standard call-by-need semantics. The standard operational semantics for the call-by-need evaluation is Launchbury’s Natural Semantics [Launchbury 1993]. Expressions are stored unevaluated in mutable heaps, which are evaluated when needed and the results are memoised by modifying existing bindings in the heaps. An extension of the semantics in the same paper is to annotate the evaluation judgements with a count of function applications for tracking computation cost. While the Natural Semantics intuitively captures both demand-drivenness and sharing, it is difficult to reason about, thus promoting various efforts in developing new semantics for analysing call-by-need programs and their computation cost.

Clairvoyance Semantics and Clairvoyance Monad. The Clairvoyance Semantics provides an alternative approach to the standard semantics by interpreting the call-by-need evaluation as call-by-value with nondeterministic choices of proceeding or skipping certain computations [Hackett and Hutton 2019]. Li et al. [2021] defined a monadic variant of the Clairvoyance Semantics and developed it into a formal framework for verifying lazy evaluation cost. They proposed an option over a monadic *thunk* datatype to model the choice of whether to evaluate or skip the evaluation of a value of that datatype.

The monad is used for encapsulating a computation that produces a value and incurring some cost (in terms of number of ticks) and a proposition specifying some property about the computation

1275 results or cost. Due to the nondeterministic nature of the semantics, the embedded program does
 1276 not compute in the usual sense, but instead leads to Rocq Prover propositions as proof obligations.
 1277 The key benefit of this monadic formalisation of the Clairvoyance Semantics is that it avoids the
 1278 exponential explosion in the cost analysis of the underlying nondeterministic semantics. However,
 1279 it also makes programs non-executable. A user of the framework generally specifies two kinds
 1280 of specifications for a program, one for all nondeterministic evaluation branches that succeed in
 1281 computation, and one for the existence of a successful branch that exhibits some more specific
 1282 properties (similar to Incorrectness Logic [O’Hearn 2020]). When verifying a specification, one
 1283 typically needs to explicitly reason about the nondeterminism presented in the semantics by
 1284 manually making the appropriate choice at the nondeterministic branches. In comparison, the
 1285 Memorist Semantics is deterministic, executable, and tracks and propagates all relevant information
 1286 alongside the computation. On the other hand, the Memorist Semantics tracks usage sets, which
 1287 can be potentially challenging for formal reasoning in the Rocq Prover.

1288 *Demand Semantics.* The more recent Demand Semantics by Xia et al. [2024] is another approach
 1289 to reason about demand and cost for lazily evaluated programs. The semantics is adapted from
 1290 the semantics of Bjerner and Holmström [1989] to a total and typed setting. It also adopts the
 1291 monadic thunk type from the Clairvoyance framework, but assigns it an additional meaning as
 1292 a representation of demand. Two kinds of demand are considered in the semantics: the output
 1293 demand, which is externally given and represents the demand on the output value in the usual
 1294 sense, and the input demand, which describes how evaluated the input needs to be for a call-by-
 1295 need evaluation driven the output demand. The semantics is mainly concerned with inferring the
 1296 minimum input demand from the output demand. To achieve this purpose, the Demand Semantics
 1297 defines a forward evaluation and a backward evaluation. Both directions are deterministic, thus
 1298 guaranteeing that backward inference of a minimal input demand is always possible for a given
 1299 output demand. However, having two separate semantics and two copies of the code is not ideal.
 1300

1301 The Demand Semantics aims at analysing the minimal demand on the input under lazy evaluation
 1302 with respect to the demand on the output. Accordingly, the backward inference is defined to always
 1303 take a specific output demand (along with the fully evaluated value) as input. For the same forward
 1304 function applied on the same input (the usual input to the forward evaluation) but with a different
 1305 output demand specified, the semantics must redo the evaluation with the new output demand.
 1306 Meanwhile, the Memorist Semantics focuses on utilising information gathered once and for all
 1307 in the analysis of subsequent computations. Analysis is based on information that is tracked and
 1308 propagated during evaluation, which does not vary according to output demand. Hence, analysis
 1309 for the same function requires no re-evaluation when different demands on outputs are specified.

1310 *Formally verifying the cost of lazy evaluation.* Handley et al. [2020] implemented a monadic
 1311 framework for analysing the computation cost of pure Haskell programs in Liquid Haskell [Vazou
 1312 2016], a proof assistant based on extending Haskell with refinement types. Nonstrictness is built
 1313 into the relevant datatypes in this framework, sharing—or memoisation—must be handled by users
 1314 explicitly, by inserting a pay construct into appropriate places of the code. This approach is inherited
 1315 from an earlier work by Danielsson [2008] for verifying time cost bounds for lazy functional data
 1316 structures in the theorem prover Agda, and is based on an amortised analysis technique by Okasaki
 1317 [1999]. Danielsson [2008]’s library employs a type-based approach by encoding and keeping track
 1318 of cost at the type level. The core element of the library is a monadic dependent *thunk* type with
 1319 embedded information about cost. Cost analysis is facilitated by Agda’s type inference system.

1320 There are quite a number of recent efforts focusing on formally reasoning about the amortized
 1321 cost of lazy functional data structures presented by Okasaki [1999]. For example, Pottier et al. [2024]
 1322 used the Iris^{\$} framework [Mével et al. 2019] to verify the banker’s queue, the physicist’s queue,
 1323

and the implicit queue in the Rocq Prover. Their approach directly reasons about mutable cells using the Iris separation logic [Spies et al. 2022], instead of relying on any alternative semantics. Xia et al. [2024] used the Demand Semantics to formally reason about the amortized cost of the banker’s queue and the implicit queue. They also formally proved that these data structures are persistent. Van Brügge [2024] used LiquidHaskell to prove the amortized cost of a simple stack, binomial heaps, and Claessen’s variant of finger trees [Claessen 2020], but these proofs are done in a non-lazy setting where there is no sharing. Our work does not focus on formal reasoning, but we would also like to explore applying the Memorist Semantics to reason about amortized cost and persistence in the future.

Analysing and inferring resource bounds for call-by-need programs. Sands [1990] describes a method to mechanically analyse the time cost bounds for call-by-need functional programs based on strictness analysis which can be applied to higher-order functions by transforming programs to include additional structures containing information about function applications and the associated cost. On the automated reasoning side, Jost et al. [2017] present a type-based analysis that automatically infers the cost of call-by-need programs using the Automatic Amortised Resource Analysis (AARA) technique introduced by Hofmann and Jost [2003]. The core idea of AARA shares much resemblance to a variant of the amortised analysis technique. The notion of prepaying cost from the amortised analysis is employed to avoid duplicating the evaluation cost of shared expressions. Moreira et al. [2020] further extend the analysis from inferring linear bounds to handle univariate polynomial bounds, based on the method introduced by Hoffmann and Hofmann [2010].

Resource analysis for non-lazy languages. An abundance of work exists on resource analysis for languages that are not lazy. As an analysis technique and system for automatically inferring resource bounds for programs, the Automatic Amortised Resource Analysis mentioned above is first introduced by Hofmann and Jost [2003] for reasoning about linearly-bounded heap space cost of first-order strict functional programs, and has subsequently gone through further extension, e.g., [Hofmann et al. 2017; Hoffmann and Hofmann 2010; Jost et al. 2010]. The analysis has been implemented in the Resource Aware ML language [Hoffmann et al. 2012]. Carboneaux et al. [2017] develop a framework based on this analysis technique for certified automatic inference of resource bounds for low-level programs. Alongside inferring bounds for programs, the framework also generates Rocq Prover proofs to verify the bounds automatically.

Some recent efforts focus on building frameworks that unify the analysis of call-by-value and call-by-name languages, though call-by-need is often not taken into account due to difficulties with modelling laziness. The λ -amor framework developed by Rajani et al. [2021] is a time complexity analysis framework that is based on amortisation and affine types that provides unified analysis for call-by-value and call-by-name languages by directly simulating the former with the latter using a monadic approach. The dependently typed logical framework **calf** for complexity analysis, introduced by Niu et al. [2022], handles the two strategies uniformly via the use of a call-by-push-value language [Levy 2004]. The framework has been used to develop cost-aware denotational semantics for functional languages [Niu and Harper 2023] and extended by Grodin et al. [2024] to also handle other computational effects by incorporating inequational reasoning into the framework’s type theory.

Another line of work follows a different approach for analysing time complexity. Namely, this approach extracts recurrence relations in terms of input size from programs, whose closed-form solution is a function expressing the cost [Cutler et al. 2020; Danner and Licata 2022; Danner et al. 2015, 2013; Kavvos et al. 2020].

On the verification side, Wang et al. [2017] introduce the TiML language with built-in support for automatic verification of time complexity bounds of programs written in this language, by allowing

1373 users to provide complexity bound specifications in type annotations which are verified via type
 1374 inference. The type system of their work is inspired by dependent types and refinement types.
 1375 McCarthy et al. [2018] implement a Rocq Prover library for verifying program time complexity by
 1376 annotating information about cost in a monadic type. Guéneau et al. [2018] formalised the big-O
 1377 notation and developed a framework based on an extension of the Separation Logic for verifying
 1378 worst-case time complexity bounds for higher-order imperative programs in Rocq Prover.

1379 7 Limitations

1381 We note that the current work has some limitations on expressivity. Presently, the Memorist
 1382 Semantics focuses on total programs with structural recursions but without general recursion.
 1383 By limiting to total programs, we can keep the semantics simple and enable it to be shallow-
 1384 embedded in Rocq Prover, whose specification language is itself total. However, we are also
 1385 interested in extending the the current work to allow reasoning about some potentially non-
 1386 terminating programs (that do terminate) and programs that work with infinite data structures. In
 1387 practice, one potential way to simulate general recursion is to use fuel to limit the number of steps.
 1388 Since our goal is to reason about computation cost, we will generally be considering programs that
 1389 halt eventually when evaluated lazily, for which there must be fuel that is sufficient.

1390 Another limitation is the lack of higher-order functions. Higher-order functions are powerful
 1391 constructs and are frequently used in functional programming. Including them in a cost analysis
 1392 framework is thus desirable, although they can also introduce quite a few complications. In the
 1393 future, we would like to address the handling of higher-order functions in the Memorist Semantics.

1394 8 Conclusion and Future Work

1395 We have presented the Memorist Semantics, a novel cost semantics for call-by-need evaluation
 1396 that tracks both usage and computation cost in a deterministic manner. Our approach advances
 1397 on prior approaches in that it enables cost analysis that is independent of demand context, avoids
 1398 code duplication, and provides fine-grained cost attribution to each individual component of a
 1399 term. Crucially, this semantics enables analyzing the cost of lazy programs while preserving the
 1400 natural deterministic compositional call-by-value evaluation structure, making it an intuitive
 1401 practical foundation for further cost-aware program reasoning. Our semantics opens the door
 1402 for future work on verifying compiler optimizations, such as dead code elimination based on our
 1403 usage sets and verified resource-bounded computation. By formalizing our model in the Rocq
 1404 Prover, we provide a rigorous basis that opens the door for the future development of cost analysis
 1405 verification frameworks for real-world lazy functional languages such as Haskell. Furthermore,
 1406 with the semantics being computable, lazy cost can be inferred directly via executing the shallowly
 1407 embedded programs, and in principle, the evaluation only needs to be performed once for analysing
 1408 the same program with various demands. Such features can potentially be helpful in areas such as
 1409 testing. Another future goal is thus to investigate the application of the semantics to property-based
 1410 testing.

1411 References

- 1412 Danel Ahman and Tarmo Uustalu. 2013. Update Monads: Cointerpreting Directed Containers. In *19th International Conference
 1413 on Types for Proofs and Programs, TYPES 2013, April 22–26, 2013, Toulouse, France (LIPIcs, Vol. 26)*, Ralph Matthes and Aleksy
 1414 Schubert (Eds.). Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 1–23. <https://doi.org/10.4230/LIPICS.TYPES.2013.1>
- 1415 Bror Bjerner and S. Holmström. 1989. A composition approach to time analysis of first order lazy functional programs. In
 1416 *Proceedings of the fourth international conference on Functional programming languages and computer architecture - FPCA
 1417 '89 (FPCA '89)*. ACM Press, 157–165. <https://doi.org/10.1145/99370.99382>
- 1418 Joachim Breitner, Antal Spector-Zabusky, Yao Li, Christine Rizkallah, John Wiegley, Joshua M. Cohen, and Stephanie
 1419 Weirich. 2021. Ready, Set, Verify! Applying hs-to-coqm to real-world Haskell code. *J. Funct. Program.* 31 (2021), e5.

- 1422 <https://doi.org/10.1017/S0956796820000283>
 1423 Quentin Carbonneaux, Jan Hoffmann, Thomas W. Reps, and Zhong Shao. 2017. Automated Resource Analysis with Coq
 1424 Proof Objects. In *Computer Aided Verification - 29th International Conference, CAV 2017, Heidelberg, Germany, July 24-28,
 1425 2017, Proceedings, Part II (Lecture Notes in Computer Science, Vol. 10427)*, Rupak Majumdar and Viktor Kuncak (Eds.).
 1426 Springer, 64–85. https://doi.org/10.1007/978-3-319-63390-9_4
- 1427 Koen Claessen. 2020. Finger trees explained anew, and slightly simplified (functional pearl). In *Proceedings of the 13th ACM
 1428 SIGPLAN International Symposium on Haskell, Haskell@ICFP 2020, Virtual Event, USA, August 7, 2020*, Tom Schrijvers
 1429 (Ed.). ACM, 31–38. <https://doi.org/10.1145/3406088.3409026>
- 1430 Joseph W. Cutler, Daniel R. Licata, and Norman Danner. 2020. Denotational recurrence extraction for amortized analysis.
 1431 *Proc. ACM Program. Lang.* 4, ICFP (2020), 97:1–97:29. <https://doi.org/10.1145/3408979>
- 1432 Nils Anders Danielsson. 2008. Lightweight semiformal time complexity analysis for purely functional data structures.
 1433 In *Proceedings of the 35th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages, POPL 2008,
 1434 San Francisco, California, USA, January 7-12, 2008*, George C. Necula and Philip Wadler (Eds.). ACM, 133–144. <https://doi.org/10.1145/1328438.1328457>
- 1435 Norman Danner and Daniel R. Licata. 2022. Denotational semantics as a foundation for cost recurrence extraction for
 1436 functional languages. *J. Funct. Program.* 32 (2022), e8. <https://doi.org/10.1017/S095679682200003X>
- 1437 Norman Danner, Daniel R. Licata, and Ramyaa. 2015. Denotational cost semantics for functional languages with inductive
 1438 types. In *Proceedings of the 20th ACM SIGPLAN International Conference on Functional Programming, ICFP 2015, Vancouver,
 1439 BC, Canada, September 1-3, 2015*, Kathleen Fisher and John H. Reppy (Eds.). ACM, 140–151. <https://doi.org/10.1145/2784731.2784749>
- 1440 Norman Danner, Jennifer Paykin, and James S. Royer. 2013. A static cost analysis for a higher-order language. In *Proceedings of
 1441 the 7th Workshop on Programming languages meets program verification, PLPV 2013, Rome, Italy, January 22, 2013*, Matthew
 1442 Might, David Van Horn, Andreas Abel, and Tim Sheard (Eds.). ACM, 25–34. <https://doi.org/10.1145/2428116.2428123>
- 1443 Harrison Grodin, Yue Niu, Jonathan Sterling, and Robert Harper. 2024. Decalf: A Directed, Effectful Cost-Aware Logical
 1444 Framework. *Proc. ACM Program. Lang.* 8, POPL (2024), 273–301. <https://doi.org/10.1145/3632852>
- 1445 Armaël Guéneau, Arthur Charguéraud, and François Pottier. 2018. A Fistful of Dollars: Formalizing Asymptotic Complexity
 1446 Claims via Deductive Program Verification. In *Programming Languages and Systems - 27th European Symposium on
 1447 Programming, ESOP 2018, Held as Part of the European Joint Conferences on Theory and Practice of Software, ETAPS 2018,
 1448 Thessaloniki, Greece, April 14-20, 2018, Proceedings (Lecture Notes in Computer Science, Vol. 10801)*, Amal Ahmed (Ed.).
 1449 Springer, 533–560. https://doi.org/10.1007/978-3-319-89884-1_19
- 1450 Jennifer Hackett and Graham Hutton. 2019. Call-by-need is clairvoyant call-by-value. *Proc. ACM Program. Lang.* 3, ICFP
 1451 (2019), 114:1–114:23. <https://doi.org/10.1145/3341718>
- 1452 Martin A. T. Handley, Niki Vazou, and Graham Hutton. 2020. Liquidate your assets: reasoning about resource usage in
 1453 liquid Haskell. *Proc. ACM Program. Lang.* 4, POPL (2020), 24:1–24:27. <https://doi.org/10.1145/3371092>
- 1454 Jan Hoffmann, Klaus Aehlig, and Martin Hofmann. 2012. Resource Aware ML. In *Computer Aided Verification - 24th
 1455 International Conference, CAV 2012, Berkeley, CA, USA, July 7-13, 2012 Proceedings (Lecture Notes in Computer Science,
 1456 Vol. 7358)*, P. Madhusudan and Sanjit A. Seshia (Eds.). Springer, 781–786. https://doi.org/10.1007/978-3-642-31424-7_64
- 1457 Jan Hoffmann, Ankush Das, and Shu-Chun Weng. 2017. Towards automatic resource bound analysis for OCaml. In *Proceedings
 1458 of the 44th ACM SIGPLAN Symposium on Principles of Programming Languages, POPL 2017, Paris, France, January 18-20,
 1459 2017*, Giuseppe Castagna and Andrew D. Gordon (Eds.). ACM, 359–373. <https://doi.org/10.1145/3009837.3009842>
- 1460 Jan Hoffmann and Martin Hofmann. 2010. Amortized Resource Analysis with Polynomial Potential. In *Programming
 1461 Languages and Systems, 19th European Symposium on Programming, ESOP 2010, Held as Part of the Joint European
 1462 Conferences on Theory and Practice of Software, ETAPS 2010, Paphos, Cyprus, March 20-28, 2010. Proceedings (Lecture Notes
 1463 in Computer Science, Vol. 6012)*, Andrew D. Gordon (Ed.). Springer, 287–306. https://doi.org/10.1007/978-3-642-11957-6_16
- 1464 Martin Hofmann and Steffen Jost. 2003. Static prediction of heap space usage for first-order functional programs. In *Conference Record of POPL 2003: The 30th SIGPLAN-SIGACT Symposium on Principles of Programming Languages, New
 1465 Orleans, Louisiana, USA, January 15-17, 2003*, Alex Aiken and Greg Morrisett (Eds.). ACM, 185–197. <https://doi.org/10.1145/604131.604148>
- 1466 Steffen Jost, Kevin Hammond, Hans-Wolfgang Loidl, and Martin Hofmann. 2010. Static determination of quantitative
 1467 resource usage for higher-order programs. In *Proceedings of the 37th ACM SIGPLAN-SIGACT Symposium on Principles of
 1468 Programming Languages, POPL 2010, Madrid, Spain, January 17-23, 2010*, Manuel V. Hermenegildo and Jens Palsberg
 1469 (Eds.). ACM, 223–236. <https://doi.org/10.1145/1706299.1706327>
- 1470 Steffen Jost, Pedro B. Vasconcelos, Mário Florido, and Kevin Hammond. 2017. Type-Based Cost Analysis for Lazy Functional
 1471 Languages. *J. Autom. Reason.* 59, 1 (2017), 87–120. <https://doi.org/10.1007/S10817-016-9398-9>
- 1472 G. A. Kavvos, Edward Morehouse, Daniel R. Licata, and Norman Danner. 2020. Recurrence extraction for functional programs
 1473 through call-by-push-value. *Proc. ACM Program. Lang.* 4, POPL (2020), 15:1–15:31. <https://doi.org/10.1145/3371083>

- 1471 John Launchbury. 1993. A Natural Semantics for Lazy Evaluation. In *Conference Record of the Twentieth Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages, Charleston, South Carolina, USA, January 1993*, Mary S. Van Deusen and Bernard Lang (Eds.). ACM Press, 144–154. <https://doi.org/10.1145/158511.158618>
- 1472 Paul Blain Levy. 2004. *Call-By-Push-Value: A Functional/Imperative Synthesis*. Semantics Structures in Computation, Vol. 2. Springer.
- 1473 Yao Li, Li-yao Xia, and Stephanie Weirich. 2021. Reasoning about the garden of forking paths. *Proc. ACM Program. Lang.* 5, ICFP (2021), 1–28. <https://doi.org/10.1145/3473585>
- 1474 Coq development team. 2024. *The Coq proof assistant*. <https://doi.org/10.5281/zenodo.14542673>
- 1475 Jay A. McCarthy, Burke Fettscher, Max S. New, Daniel Feltey, and Robert Bruce Findler. 2018. A Coq library for internal verification of running-times. *Sci. Comput. Program.* 164 (2018), 49–65. <https://doi.org/10.1016/J.SCICO.2017.05.001>
- 1476 Glen Mével, Jacques-Henri Jourdan, and François Pottier. 2019. Time Credits and Time Receipts in Iris. In *Programming Languages and Systems - 28th European Symposium on Programming, ESOP 2019, Held as Part of the European Joint Conferences on Theory and Practice of Software, ETAPS 2019, Prague, Czech Republic, April 6–11, 2019, Proceedings (Lecture Notes in Computer Science, Vol. 11423)*, Luís Caires (Ed.). Springer, 3–29. https://doi.org/10.1007/978-3-030-17184-1_1
- 1477 Sara Moreira, Pedro B. Vasconcelos, and Mário Florido. 2020. Resource Analysis for Lazy Evaluation with Polynomial Potential. In *IFL 2020: 32nd Symposium on Implementation and Application of Functional Languages, Virtual Event / Canterbury, UK, September 2–4, 2020*, Olaf Chitil (Ed.). ACM, 104–114. <https://doi.org/10.1145/3462172.3462196>
- 1478 Yue Niu and Robert Harper. 2023. A Metalanguage for Cost-Aware Denotational Semantics. In *38th Annual ACM/IEEE Symposium on Logic in Computer Science, LICS 2023, Boston, MA, USA, June 26–29, 2023*. IEEE, 1–14. <https://doi.org/10.1109/LICS56636.2023.10175777>
- 1479 Yue Niu, Jonathan Sterling, Harrison Grodin, and Robert Harper. 2022. A cost-aware logical framework. *Proc. ACM Program. Lang.* 6, POPL (2022), 1–31. <https://doi.org/10.1145/3498670>
- 1480 Peter W. O’Hearn. 2020. Incorrectness logic. *Proc. ACM Program. Lang.* 4, POPL (2020), 10:1–10:32. <https://doi.org/10.1145/3371078>
- 1481 Chris Okasaki. 1999. *Purely functional data structures*. Cambridge University Press.
- 1482 François Pottier, Arnaël Guéneau, Jacques-Henri Jourdan, and Glen Mével. 2024. Thunks and Debts in Separation Logic with Time Credits. *Proc. ACM Program. Lang.* 8, POPL (2024). <https://hal.science/hal-04238691/file/main.pdf>
- 1483 Vineet Rajani, Marco Gaboardi, Deepak Garg, and Jan Hoffmann. 2021. A unifying type-theory for higher-order (amortized) cost analysis. *Proc. ACM Program. Lang.* 5, POPL (2021), 1–28. <https://doi.org/10.1145/3434308>
- 1484 David Sands. 1990. Complexity Analysis for a Lazy Higher-Order Language. In *ESOP’90, 3rd European Symposium on Programming, Copenhagen, Denmark, May 15–18, 1990, Proceedings (Lecture Notes in Computer Science, Vol. 432)*, Neil D. Jones (Ed.). Springer, 361–376. https://doi.org/10.1007/3-540-52592-0_74
- 1485 Antal Spector-Zabrusky, Joachim Breitner, Christine Rizkallah, and Stephanie Weirich. 2018. Total Haskell is reasonable Coq. In *Proceedings of the 7th ACM SIGPLAN International Conference on Certified Programs and Proofs, CPP 2018, Los Angeles, CA, USA, January 8–9, 2018*, June Andronick and Amy P. Felty (Eds.). ACM, 14–27. <https://doi.org/10.1145/3167092>
- 1486 Simon Spies, Lennard Gähler, Joseph Tassarotti, Ralf Jung, Robbert Krebbers, Lars Birkefeld, and Derek Dreyer. 2022. Later credits: resourceful reasoning for the later modality. *Proc. ACM Program. Lang.* 6, ICFP (2022), 283–311. <https://doi.org/10.1145/3547631>
- 1487 Jan van Brügge. 2024. Liquid Amortization: Proving Amortized Complexity with LiquidHaskell (Functional Pearl). In *Proceedings of the 17th ACM SIGPLAN International Haskell Symposium, Haskell 2024, Milan, Italy, September 6–7, 2024*, Niki Vazou and J. Garrett Morris (Eds.). ACM, 97–108. <https://doi.org/10.1145/3677999.3678282>
- 1488 Niki Vazou. 2016. *Liquid Haskell: Haskell as a Theorem Prover*. Ph. D. Dissertation. University of California, San Diego, USA. <http://www.escholarship.org/uc/item/8dm057ws>
- 1489 Peng Wang, Di Wang, and Adam Chlipala. 2017. TiML: a functional language for practical complexity analysis with invariants. *Proc. ACM Program. Lang.* 1, OOPSLA (2017), 79:1–79:26. <https://doi.org/10.1145/3133903>
- 1490 Li-yao Xia, Laura Israel, Maite Kramarz, Nicholas Coltharp, Koen Claessen, Stephanie Weirich, and Yao Li. 2024. Story of Your Lazy Function’s Life: A Bidirectional Demand Semantics for Mechanized Cost Analysis of Lazy Programs. *Proc. ACM Program. Lang.* 8, ICFP (2024), 30–63. <https://doi.org/10.1145/3674626>
- 1491
1492
1493
1494
1495
1496
1497
1498
1499
1500
1501
1502
1503
1504
1505
1506
1507
1508
1509
1510
1511
1512
1513
1514
1515
1516
1517
1518
1519